

POSTGRES MVCC

**A DEVELOPER CENTRIC VIEW OF
MULTI VERSION CONCURRENCY CONTROL**

By: Robert Sosinski

Reactive.IO

Robert Sosinski

Founder & Engineering Fellow

AGENDA

MVCC: what it is and it matters

Transactions: more than just an undo button

Isolation Levels: seeing what you need to see

Locking: control when your data is written

Cursors: stream chronologically correct data

Summary: bringing it all together

Questions: ready, fire, aim

WHAT IS MVCC?

Multi

Version

Concurrency

Control

WHAT IS MVCC?

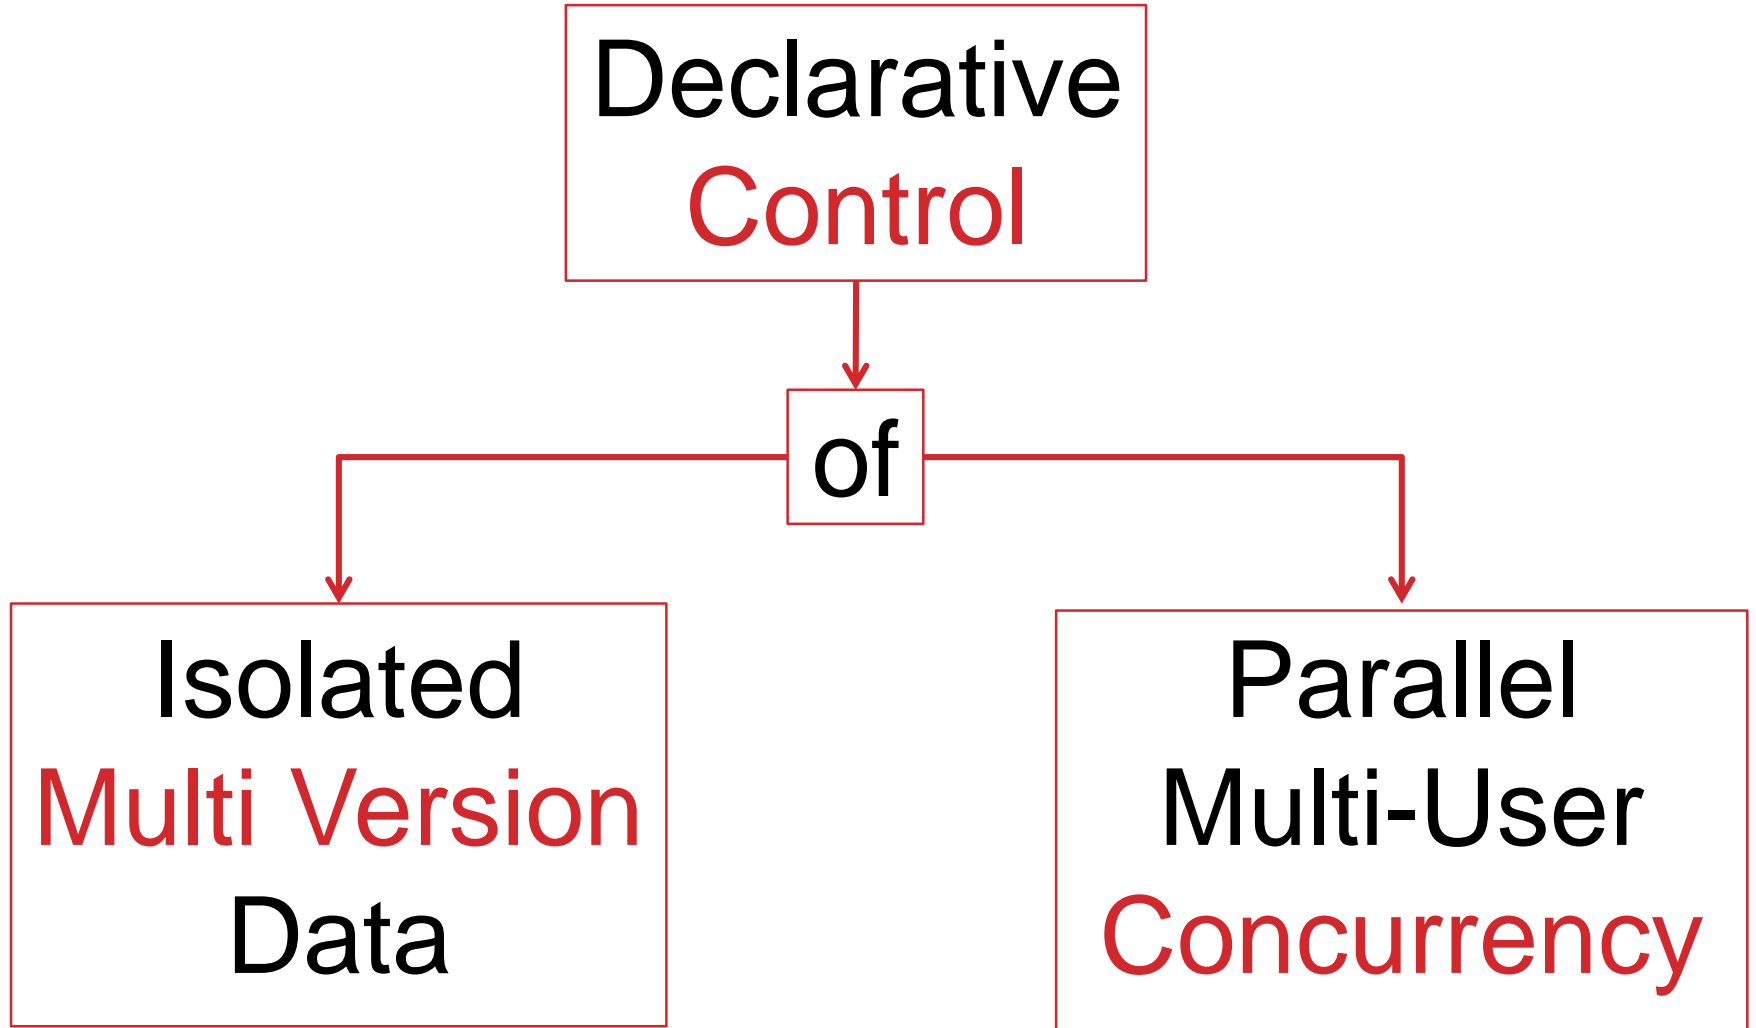
Multi

Version

Concurrency

Control

WHAT IS MVCC USING BOXES AND ARROWS



WHAT DOES MVCC DO?

- 1: Multiple users are able access the same data at the same time.
- 2: Every user sees their own isolated snapshot of the database.
- 3: Changes made by one user, will not be seen by any other user until their transaction is committed.

THE MVCC SALES PITCH

MULTI VERSION

Atomic Updates

Consistent Data

Isolated Reads

CONCURRENCY

Higher efficiency

Simpler operations

Engineering agility

A MVCC WORLD



A WORLD WITHOUT MVCC

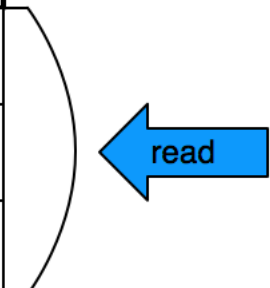
**Simple HR database
for a fictitious company
with high employee churn
without MVCC.**

UPDATE IN PLACE 1: TABLE


id	name	notes
1	Alice	Great at programming
2	Bob	Always talking to alice
3	Eve	Listens to everyone's conversations

UPDATE IN PLACE 2: SCAN

id	name	notes
1	Alice	Great at programming
2	Bob	Always talking to alice
3	Eve	Listens to everyone's conversations




UPDATE IN PLACE 3: UPDATE



id	name	notes
1	Alice	Great at programming
2	Bob	Always talking to alice
3	Eve	Listens to everyone's conversations

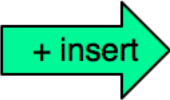
UPDATE IN PLACE 4: UPDATED



id	name	notes
1	Alice	Great at programming
2	Bob	Working very hard
3	Eve	Listens to everyone's conversations

UPDATE IN PLACE 5: INSERT

id	name	notes
1	Alice	Great at programming
2	Bob	Working very hard
3	Eve	Listens to everyone's conversations
4	Dave	Very promising new-hire



UPDATE IN PLACE 6: DELETE



id	name	notes
1	Alice	Great at programming
2	Bob	Working very hard
3	Eve	Listens to everyone's conversations
4	Dave	Very promising new-hire

UPDATE IN PLACE: 7

id	name	notes
1	Alice	Great at programming
2	Bob	Working very hard
4	Dave	Very promising new-hire

UPDATE IN PLACE 8: REALITY

	id	name	notes	
	1	Alice	Great at programming	
~ update	2	Bob	Works very hard	← read
- delete	3	<i>Eve</i>	<i>Listens to everyone's conversations</i>	
+ insert	4	Dave	Very promising new-hire	

HOW TO SOLVE THIS PROBLEM

Pessimistic locking: lock everything during writes

Imperative controls: synchronization and mutexes

System build out: everyone gets their own database

Let the cards fall: whatever happens, happens...

HOW TO SOLVE THIS PROBLEM

Pessimistic locking: lock everything during writes

Imperative controls: synchronization and mutexes

System build out: everyone gets their own database

Let the cards fall: whatever happens, happens


MVCC: Let the database handle the particulars

MVCC 1: TABLE

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
101	0	2	Bob	Always talking to alice
102	0	3	Eve	Listens to everyone's conversations


TXID
103

MVCC 2: UPDATE

	xmin	xmax	id	name	notes
	100	0	1	Alice	Great at programming
	101	0	2	Bob	Always talking to alice
	102	0	3	Eve	Listens to everyone's conversations

TXID
103

MVCC 3: UPDATE IN PROGRESS

	xmin	xmax	id	name	notes
	100	0	1	Alice	Great at programming
	101	103	2	Bob	Always talking to alice
	102	0	3	Eve	Listens to everyone's conversations

TXID
103

MVCC 4: UPDATE IN PROGRESS

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
101	103	2	Bob	Always talking to alice
102	0	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard

+ update

TXID
103

MVCC 5: UPDATED

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
101	103	2	Bob	Always talking to alice
102	0	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard

TXID
104

MVCC 6: INSERT


xmin	xmax	id	name	notes	TXID
100	0	1	Alice	Great at programming	
101	103	2	Bob	Always talking to alice	
102	0	3	Eve	Listens to everyone's conversations	
103	0	2	Bob	Working very hard	
+ insert 104	0	4	Dave	Very promising new-hire	104

MVCC 7: INSERTED

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
101	103	2	Bob	Always talking to alice
102	0	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire

TXID
105

MVCC 8: DELETE

	xmin	xmax	id	name	notes
	100	0	1	Alice	Great at programming
	101	103	2	Bob	Always talking to alice
	102	105	3	Eve	<i>Listens to everyone's conversations</i>
	103	0	2	Bob	Working very hard
	104	0	4	Dave	Very promising new-hire

TXID

105

MVCC 9: DELETED

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
101	103	2	Bob	Always talking to alice
102	105	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire

TXID
106

LETS CONCURRENTLY WRITE SOME DATA

Open Postgres Terminal

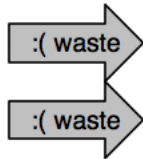
VACUUM 1: UIP TABLE

id	name	notes
1	Alice	Great at programming
2	Bob	Working very hard
4	Dave	Very promising new-hire

VACUUM 2: MVCC TABLE

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
101	103	2	Bob	Always talking to alice
102	105	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire

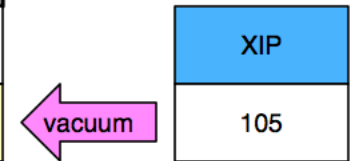
VACUUM 3: WASTE



xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
101	103	2	Bob	Always talking to alice
102	105	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire

VACUUM 4: FIRST VACUUM

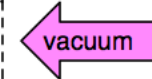
xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
101	103	2	Bob	Always talking to alice
102	105	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire



VACUUM 5: FIRST VACUUM

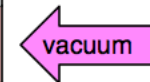
xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
102	105	3	Eve	<i>Listens to everyone's conversations</i>
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire

XIP
105



VACUUM 6: FIRST VACUUM

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
102	105	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire



XIP
105

VACUUM 7: VACUUM FINISHED

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
102	105	3	Eve	Listens to everyone's conversations
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire

XIP
105

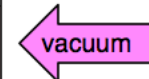
VACUUM 8: SECOND VACUUM

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
102	105	3	Eve	<i>Listens to everyone's conversations</i>
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire

XIP

VACUUM 9: SECOND VACUUM

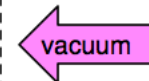
xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
102	105	3	Eve	<i>Listens to everyone's conversations</i>
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire



XIP

VACUUM 10: SECOND VACUUM

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire



XIP

VACUUM 11: VACUUM FINISHED

xmin	xmax	id	name	notes
100	0	1	Alice	Great at programming
103	0	2	Bob	Working very hard
104	0	4	Dave	Very promising new-hire

XIP

ISOLATION LEVELS

Level	Dirty Read	Nonrepeatable Read	Phantom Read	Serialization Anomaly
Read Committed <i>Default</i>	✓			
Repeatable Read	✓	✓	✓	
Serializable	✓	✓	✓	✓

Referenced from: <http://www.postgresql.org/docs/9.3/static/transaction-iso.html>

✓ Not Possible

LETS ISOLATE SOME TRANSACTIONS

Open Postgres Terminal

EXPLICIT LOCKING

TABLE LEVEL

Very broad

8 types

Can affect querying

ROW LEVEL

Very granular

2 types

Will not affect querying

ROW LEVEL LOCKING

Name	Lock Type	Blocks Update	Blocks Select For Update
For Share	Row Share	✓	
select * from people where id = 1 <u>for share</u> ;			
For Update	Row Exclusive	✓	✓
select * from people where id = 1 <u>for update</u> ;			

LETS LOCK SOME ROWS

Open Postgres Terminal

CURSORS

Streaming: break large datasets in smaller segments

Efficient: reduce a queries memory consumption

Isolated: return chronologically correct data

Traversable: can scan forward, backwards and more

Flexible: PL/pgSQL functions can return/accept cursors

LETS USE A CURSOR

Open Postgres Terminal

SUMMARY

Powerful: interact with your data on your terms

Declarative: easy to use, less chance of mistakes

Efficient: use less resources to work with more data

Scalable: handle more processes with larger volume

Flexible: do what you need do when you need it

CONDENSED SUMMARY

Control of your

Concurrent

Multi Versioned

Data

CONDENSED SUMMARY

Control of your

Concurrent

Multi Versioned

Business

THANKS

Open For Questions