Redis Sentinel

charsyam@naver.com

Redis.io

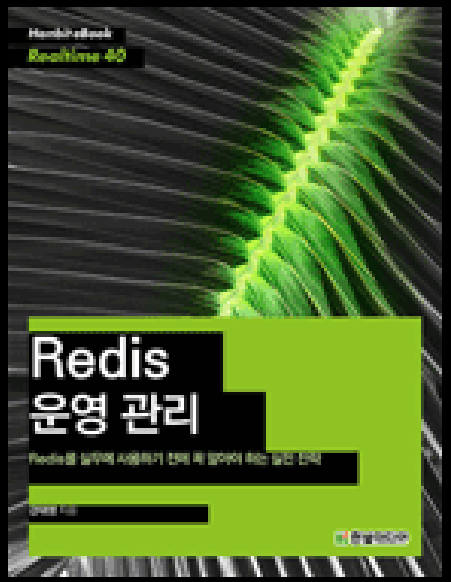

Redis/Twemproxy
contributor

카카오 홈 개발자

Redis Monitoring?

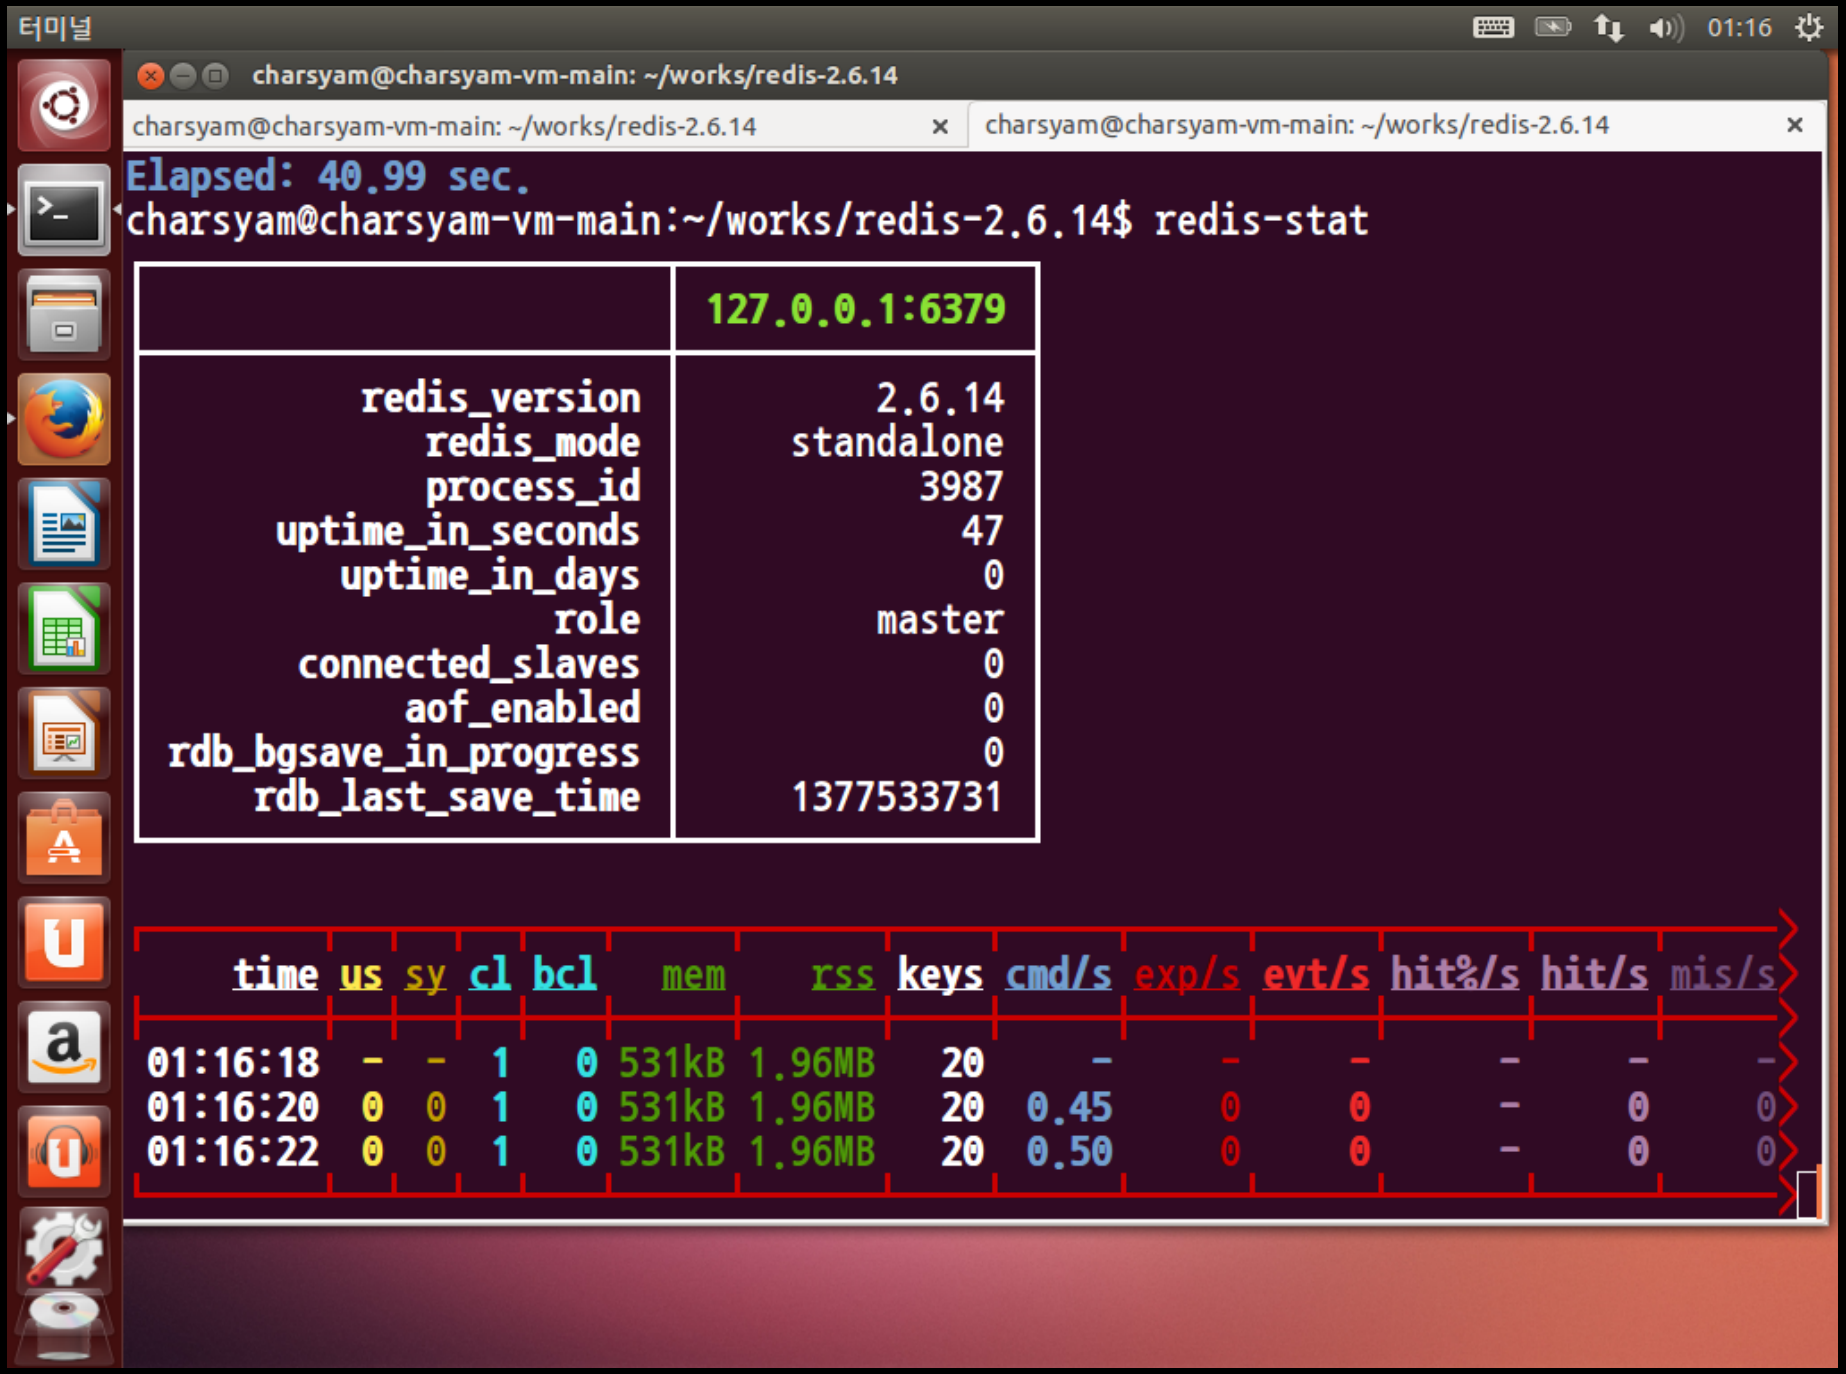
터미널
01:16
charsyam@charsyam-vm-main: ~/works/redis-2.6.14
charsyam@charsyam-vm-main: ~/works/redis-2.6.14
charsyam@charsyam-vm-main: ~/works/redis-2.6.14
Elapsed: 40.99 sec.
charsyam@charsyam-vm-main:~/works/redis-2.6.14$ redis-stat
127.0.0.1:6379
redis_version 2.6.14
redis_mode standalone
process_id 3987
uptime_in_seconds 47
uptime_in_days 0
role master
connected_slaves 0
aof_enabled 0
rdb_bgsave_in_progress 0
rdb_last_save_time 1377533731
time us sy cl bcl mem rss keys cmd/s exp/s evt/s hit%/s hit/s mis/s
01:16:18 - - 1 0 531kB 1.96MB 20 - - - - - -
01:16:20 0 0 1 0 531kB 1.96MB 20 0.45 0 0 - 0 0
01:16:22 0 0 1 0 531kB 1.96MB 20 0.50 0 0 - 0 0

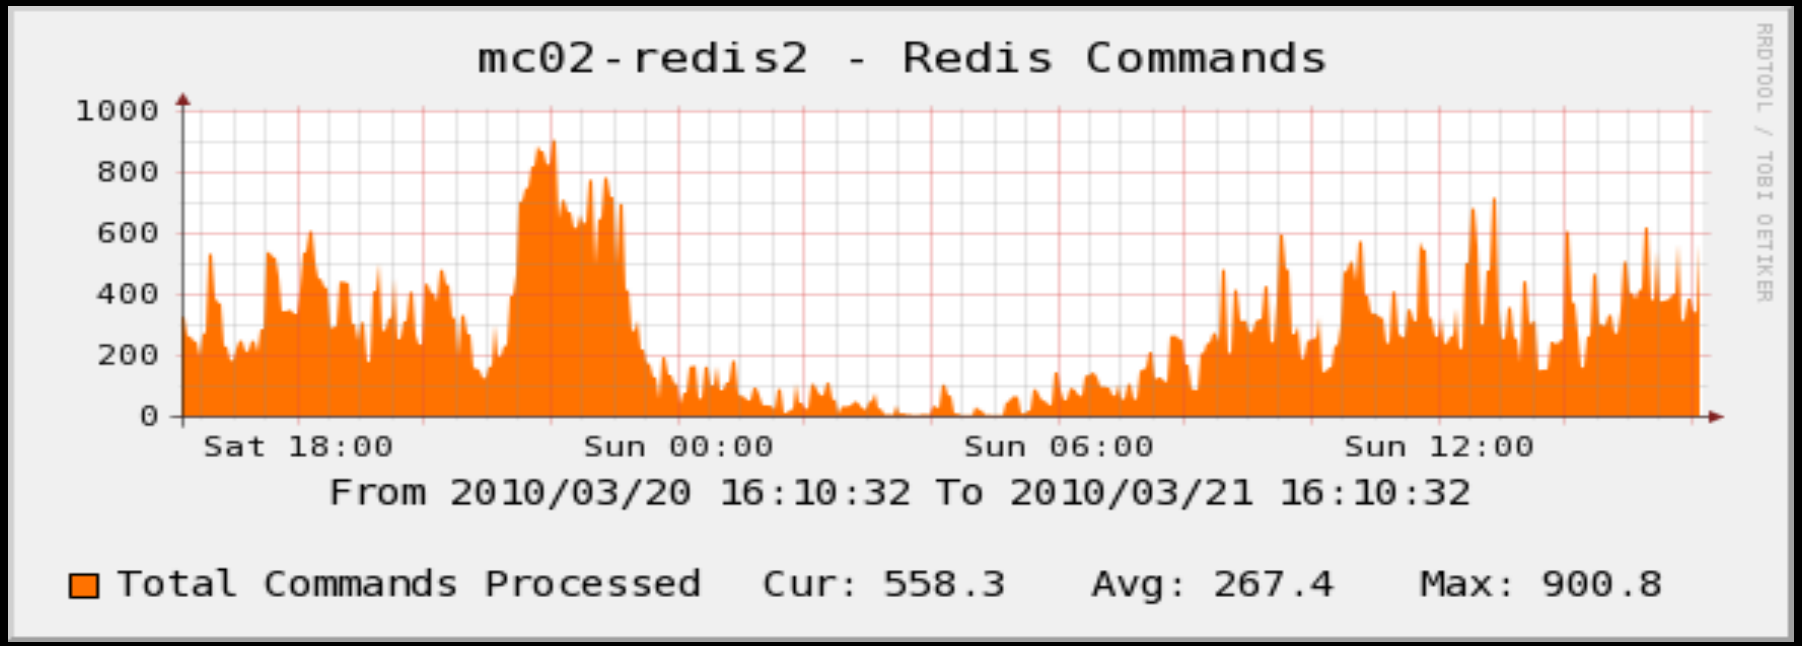
mc02-redis2 - Redis Commands
1000
800
600
400
200
0
Sat 18:00
Sun 00:00
Sun 06:00
Sun 12:00
From 2010/03/20 16:10:32 To 2010/03/21 16:10:32
Total Commands Processed Cur: 558.3 Avg: 267.4 Max: 900.8
RRDTOOL / TOBI OETIKER

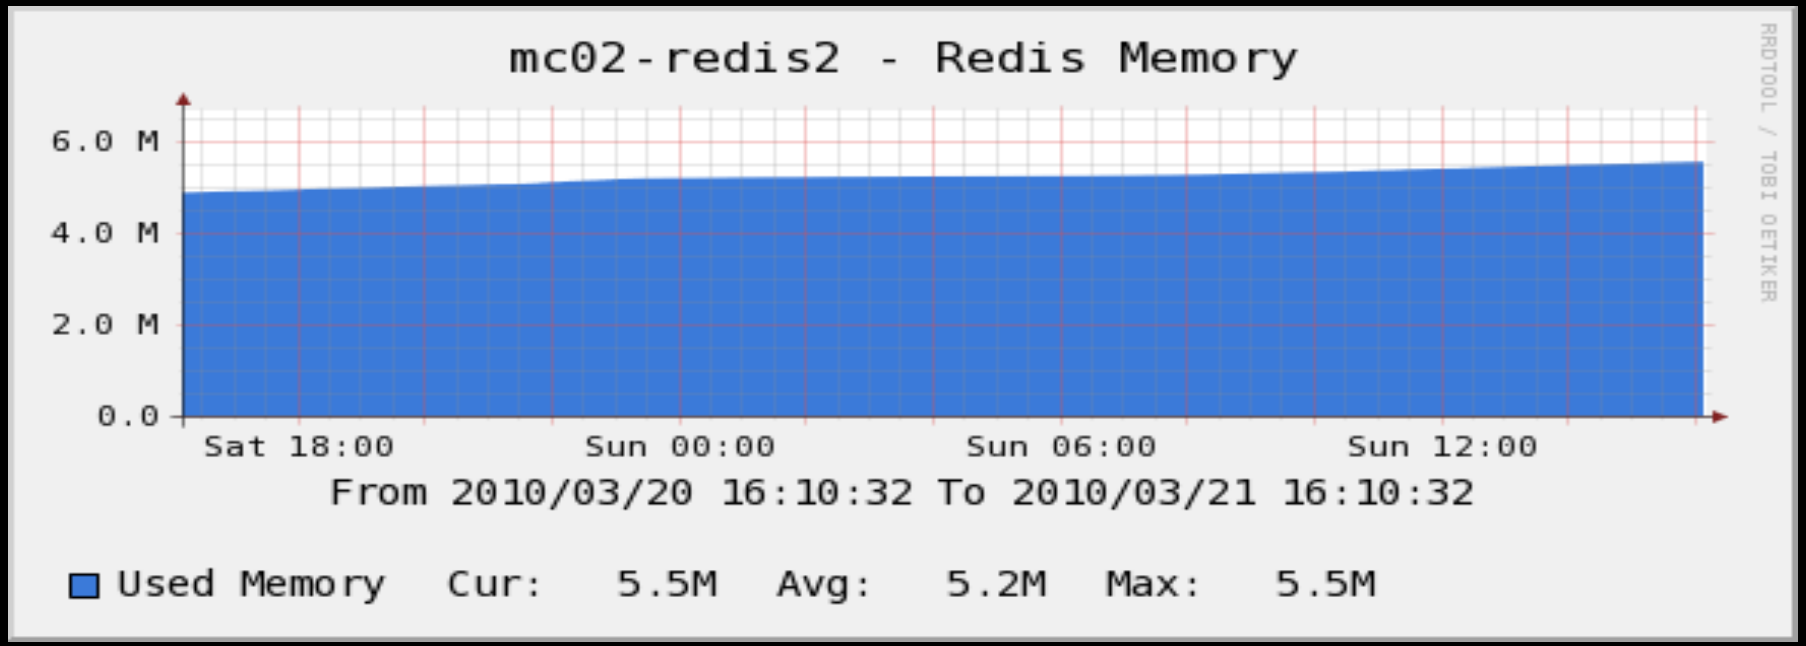
mc02-redis2 - Redis Memory
6.0 M
4.0 M
2.0 M
0.0
Sat 18:00
Sun 00:00
Sun 06:00
Sun 12:00
From 2010/03/20 16:10:32 To 2010/03/21 16:10:32
Used Memory Cur: 5.5M Avg: 5.2M Max: 5.5M
RRDTOOL / TOBI OETIKER

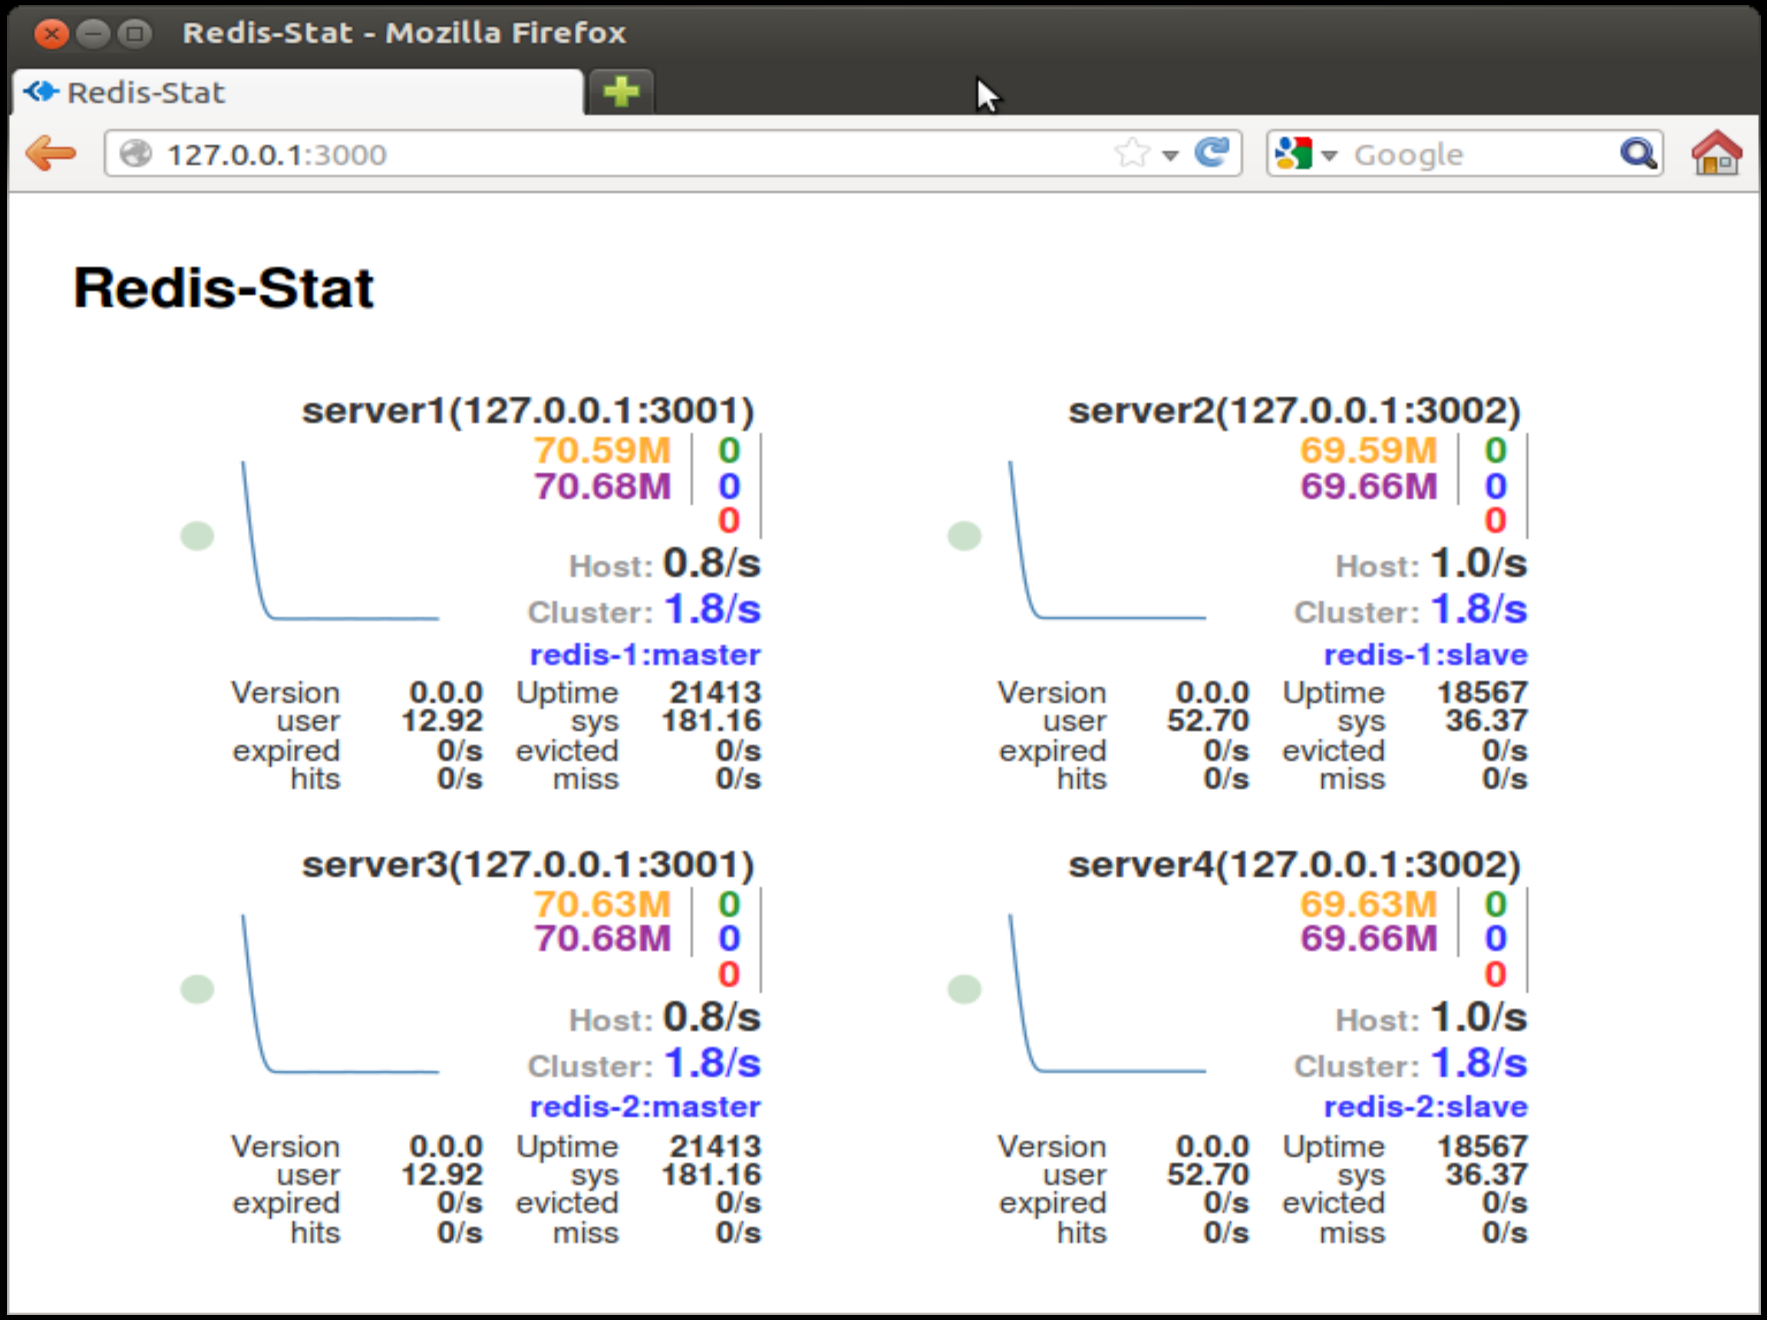
Redis-Stat - Mozilla Firefox
Redis-Stat
127.0.0.1:3000
Google
Redis-Stat
server1(127.0.0.1:3001)
70.59M
70.68M
0
0
0
Host: 0.8/s
Cluster: 1.8/s
redis-1:master
Version 0.0.0 Uptime 21413
user 12.92 sys 181.16
expired 0/s evicted 0/s
hits 0/s miss 0/s
server2(127.0.0.1:3002)
69.59M
69.66M
0
0
0
Host: 1.0/s
Cluster: 1.8/s
redis-1:slave
Version 0.0.0 Uptime 18567
user 52.70 sys 36.37
expired 0/s evicted 0/s
hits 0/s miss 0/s
server3(127.0.0.1:3001)
70.63M
70.68M
0
0
0
Host: 0.8/s
Cluster: 1.8/s
redis-2:master
Version 0.0.0 Uptime 21413
user 12.92 sys 181.16
expired 0/s evicted 0/s
hits 0/s miss 0/s
server4(127.0.0.1:3002)
69.63M
69.66M
0
0
0
Host: 1.0/s
Cluster: 1.8/s
redis-2:slave
Version 0.0.0 Uptime 18567
user 52.70 sys 36.37
expired 0/s evicted 0/s
hits 0/s miss 0/s

Checking Liveness

Checking Liveness

```
servers = [ ('localhost', 6379),
            ('localhost', 6380),
            ('localhost', 6381) ]


conns = []


def check():
    for conn in conns:
        print conn.ping()
```

checking it works

Checking it works

```
def check():
    for conn in conns:
        i = conn.info()
        print i['run_id'],
              i['used_memory_human'],
              i['total_commands_processed']
```

Is it Reliable?

Is Server really Failed?

Split Network.

Server is busy.

Monitoring is not the one we really want.

HA

High Availity

what do you do?
when your servers failed

1. Recognize Redis is down.

2. choose new master candidate

3. Promote slave to Master.

4. make client connect to new master

what Redis Sentinel do?

1. Check Redis Liveness

1. Check Redis Liveness

-Subjective Down

-Objective Down

2. choose good slave.

Good Slave

- NOT SDOWN, ODOWN, DISCONNECTED
- NOT DEMOTE
- Ping reply > info_validity_time
- Slave_priority != 0
- Info reply > info_validity_time

choose new master

- Sort
- Slave Priority
- Runid

3. Promote it

Promoting

- Send Slaveof no one to new master
- Send Slaveof [new master ip] [addr] to other redis
- Notify new master to clietns +switch-master
- Set DEMOTE mark to old-Master

Failover Scenario

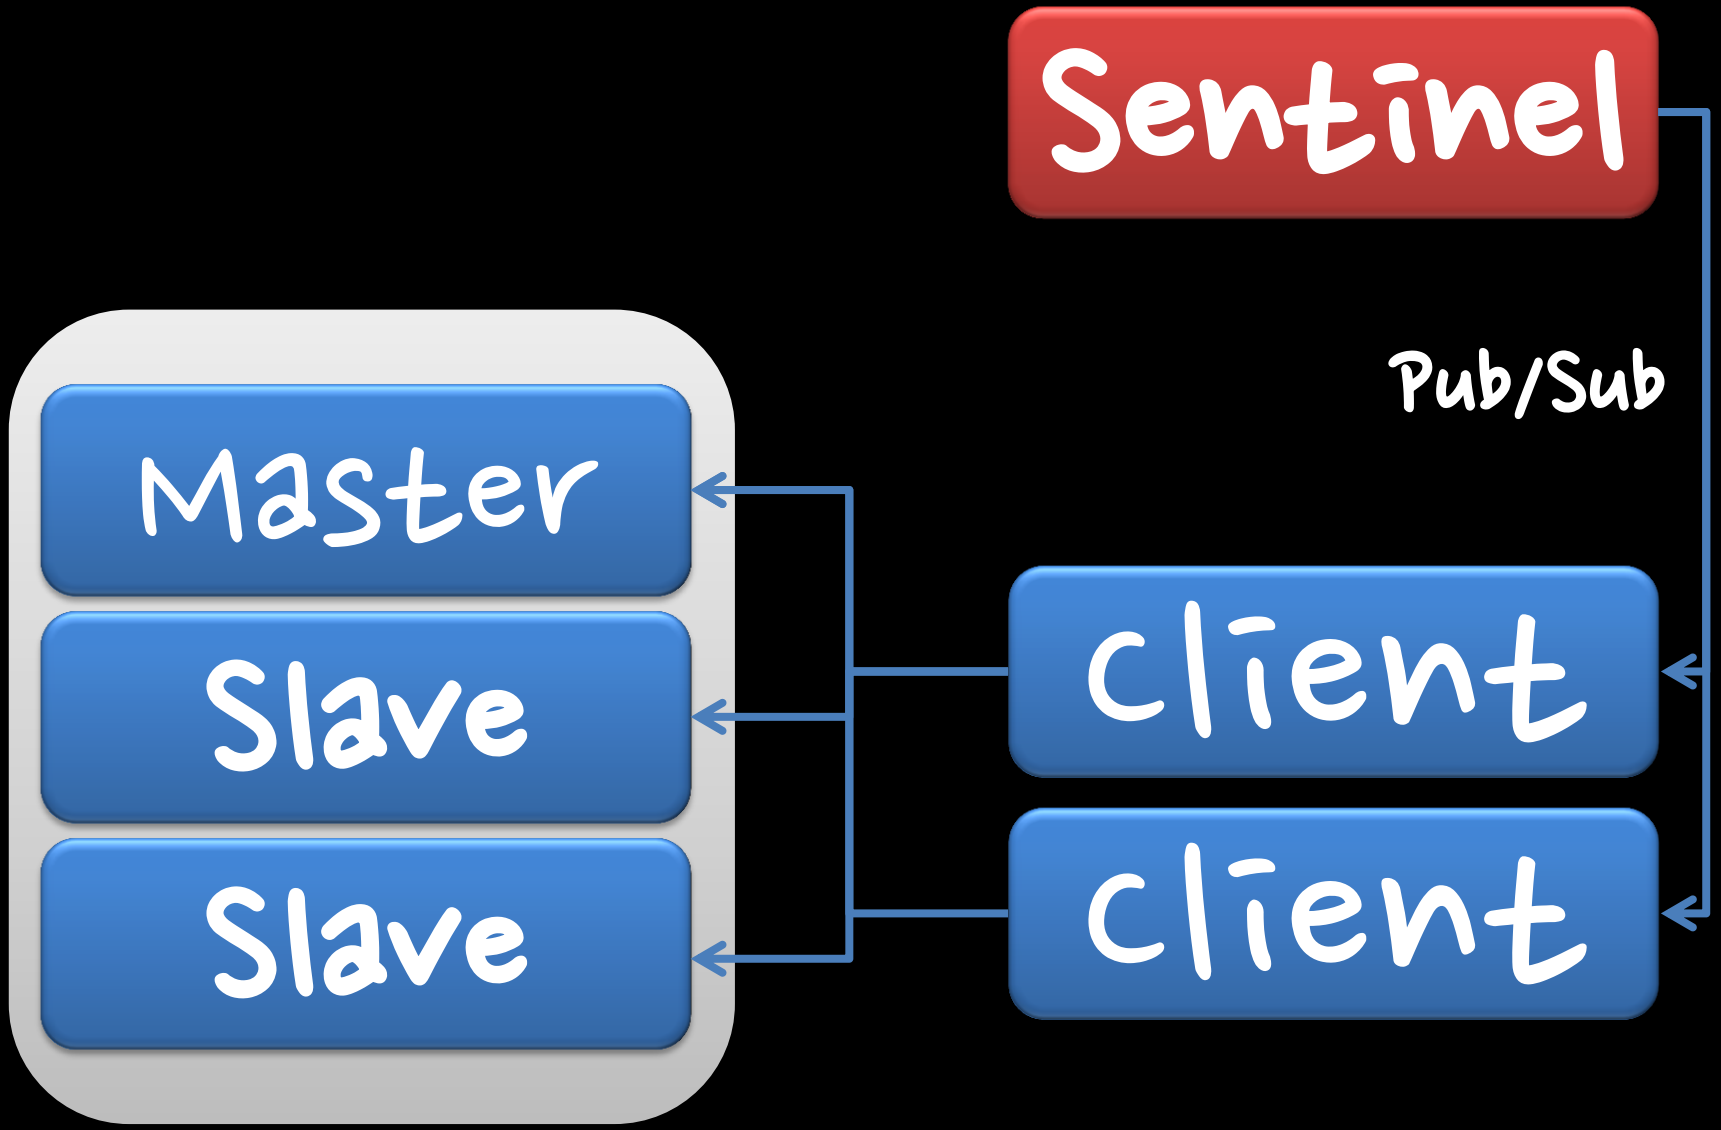
Sentinel
Pub/Sub
Master
Slave
Slave
client
client

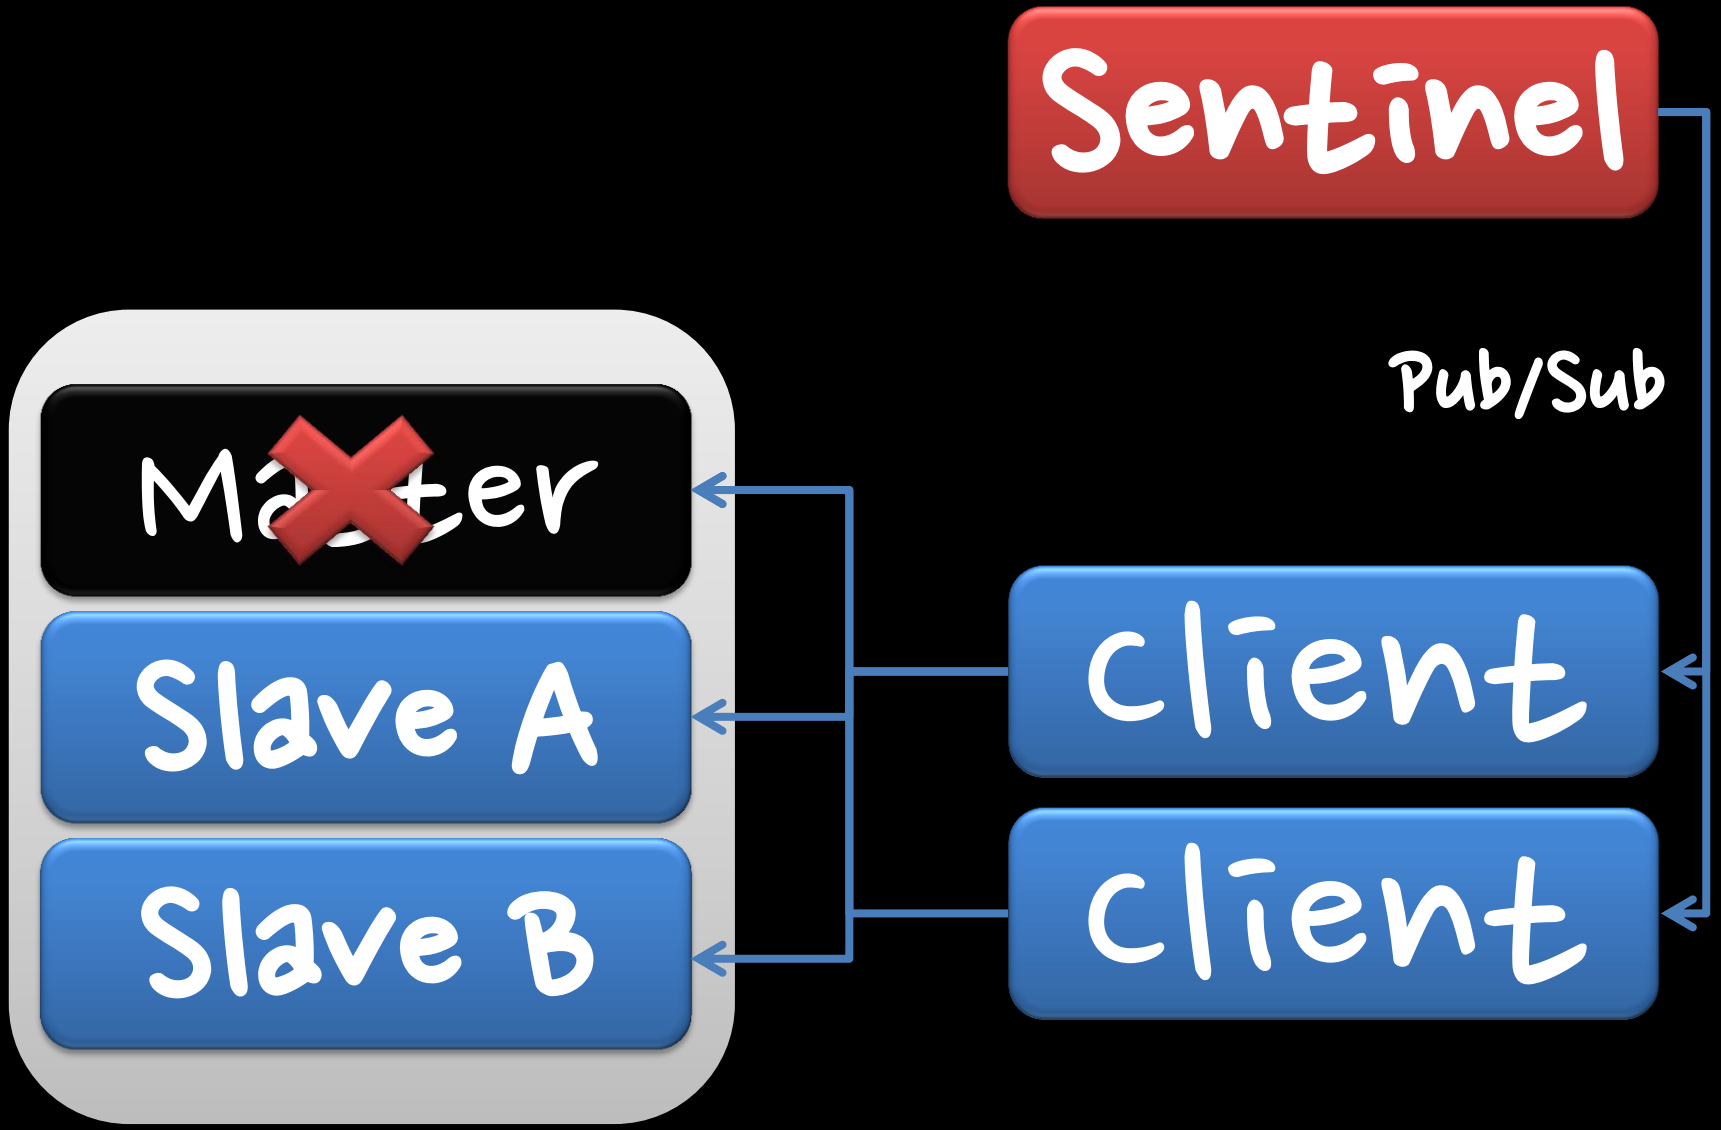
Sentinel
Pub/Sub
Master
Slave A
Slave B
client
client

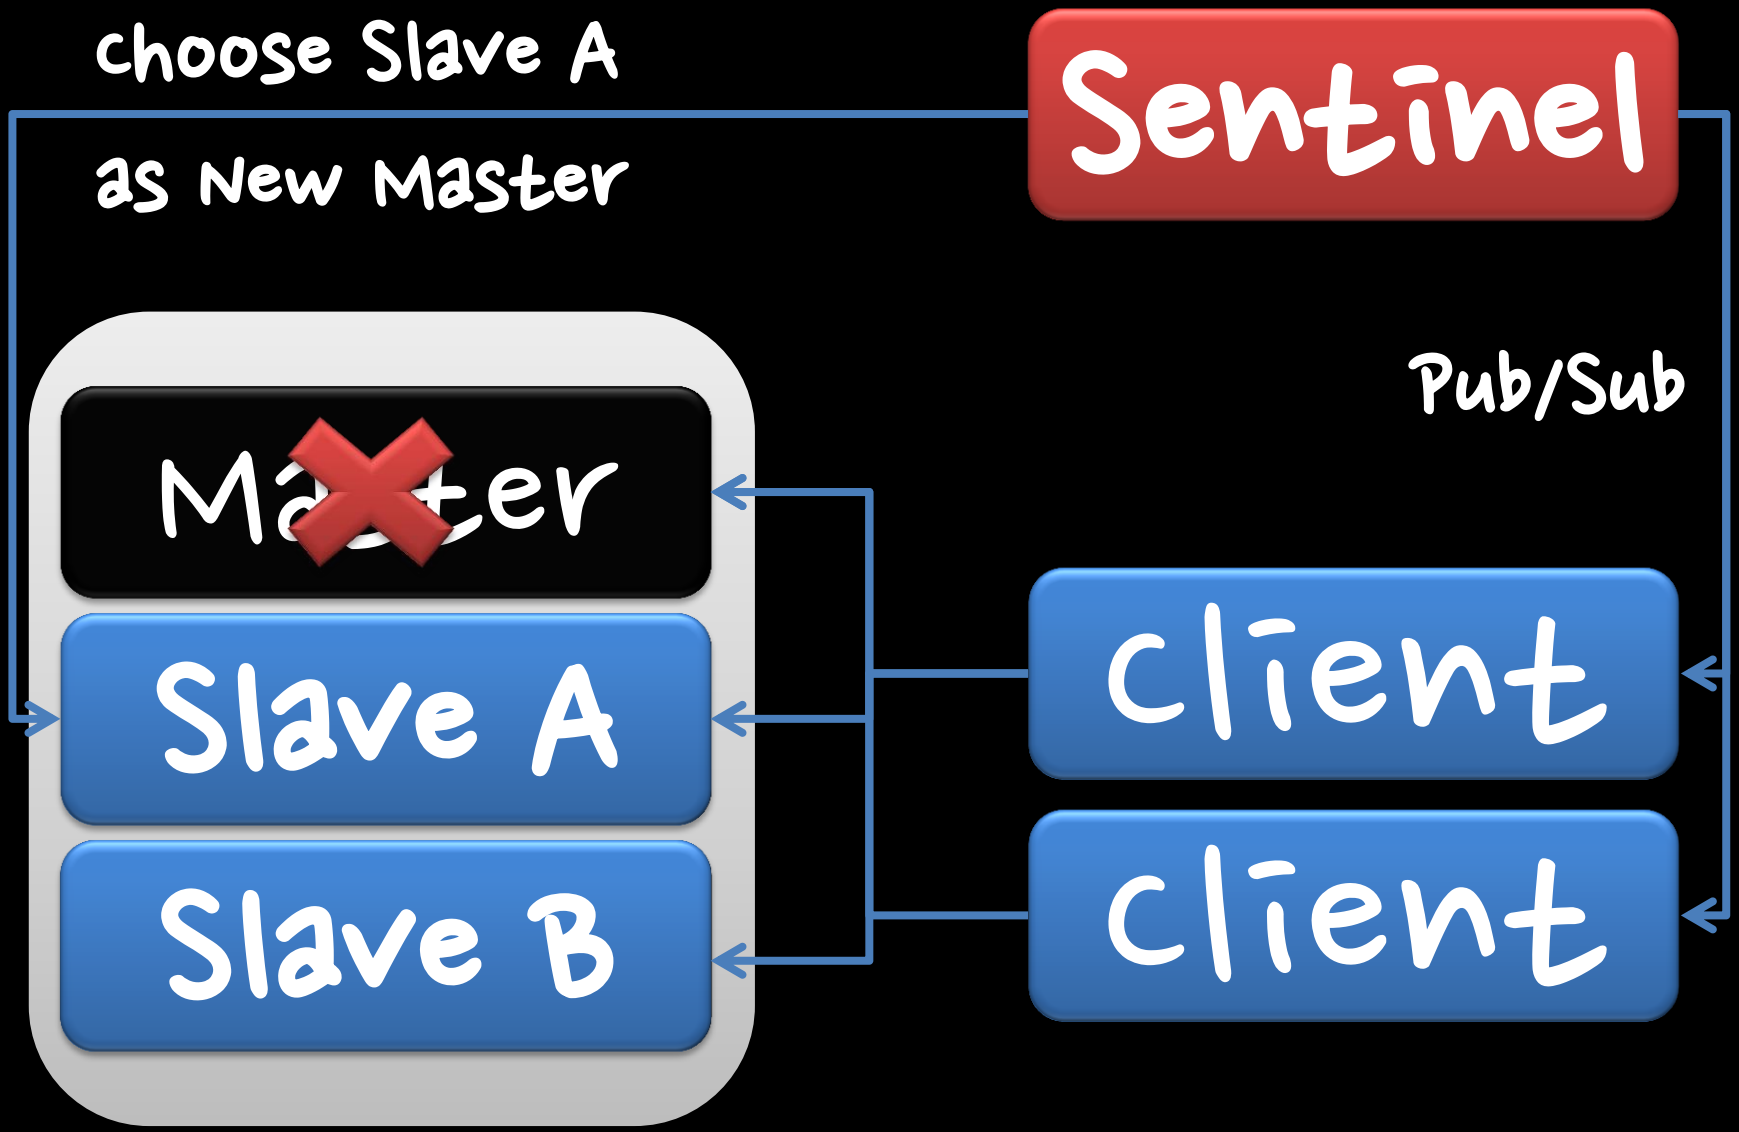
choose Slave A
as New Master
Sentinel
Pub/Sub
Master
Slave A
Slave B
client
client

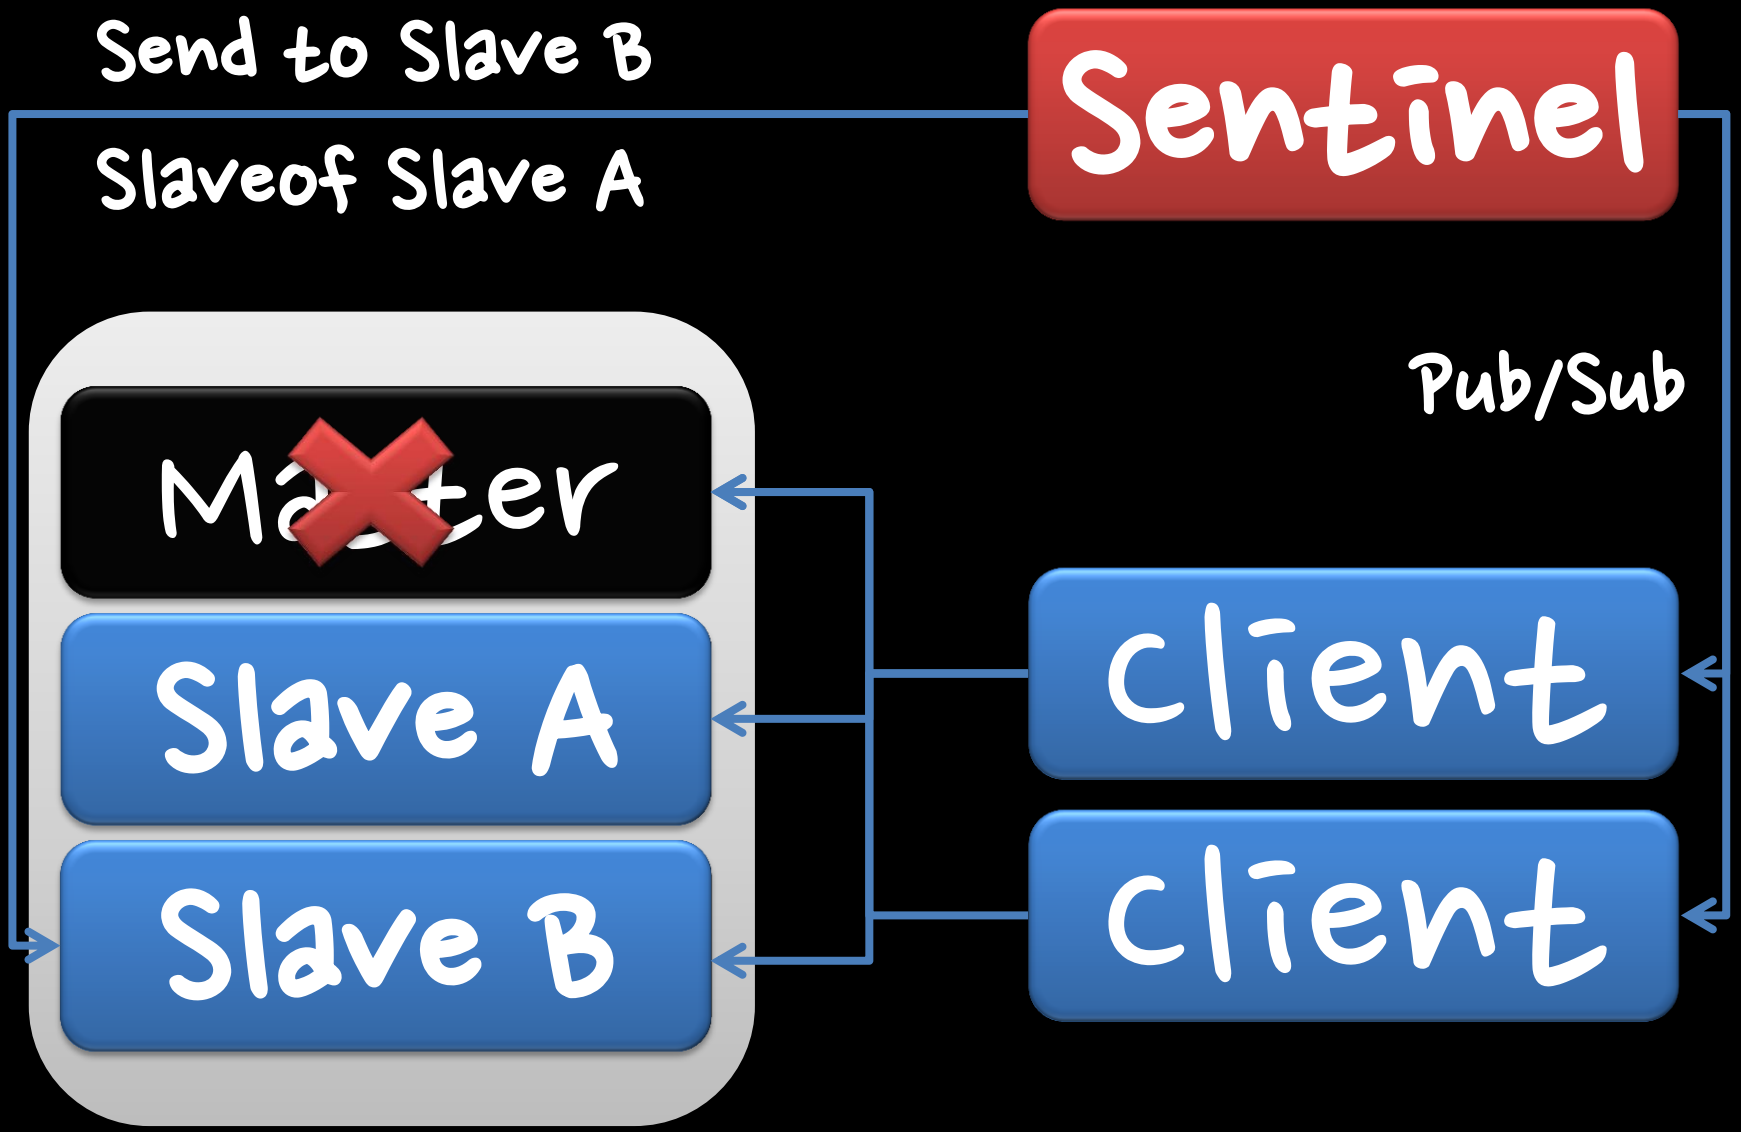
Send to Slave B
Slaveof Slave A
Sentinel
Pub/Sub
Master
Slave A
Slave B
client
client

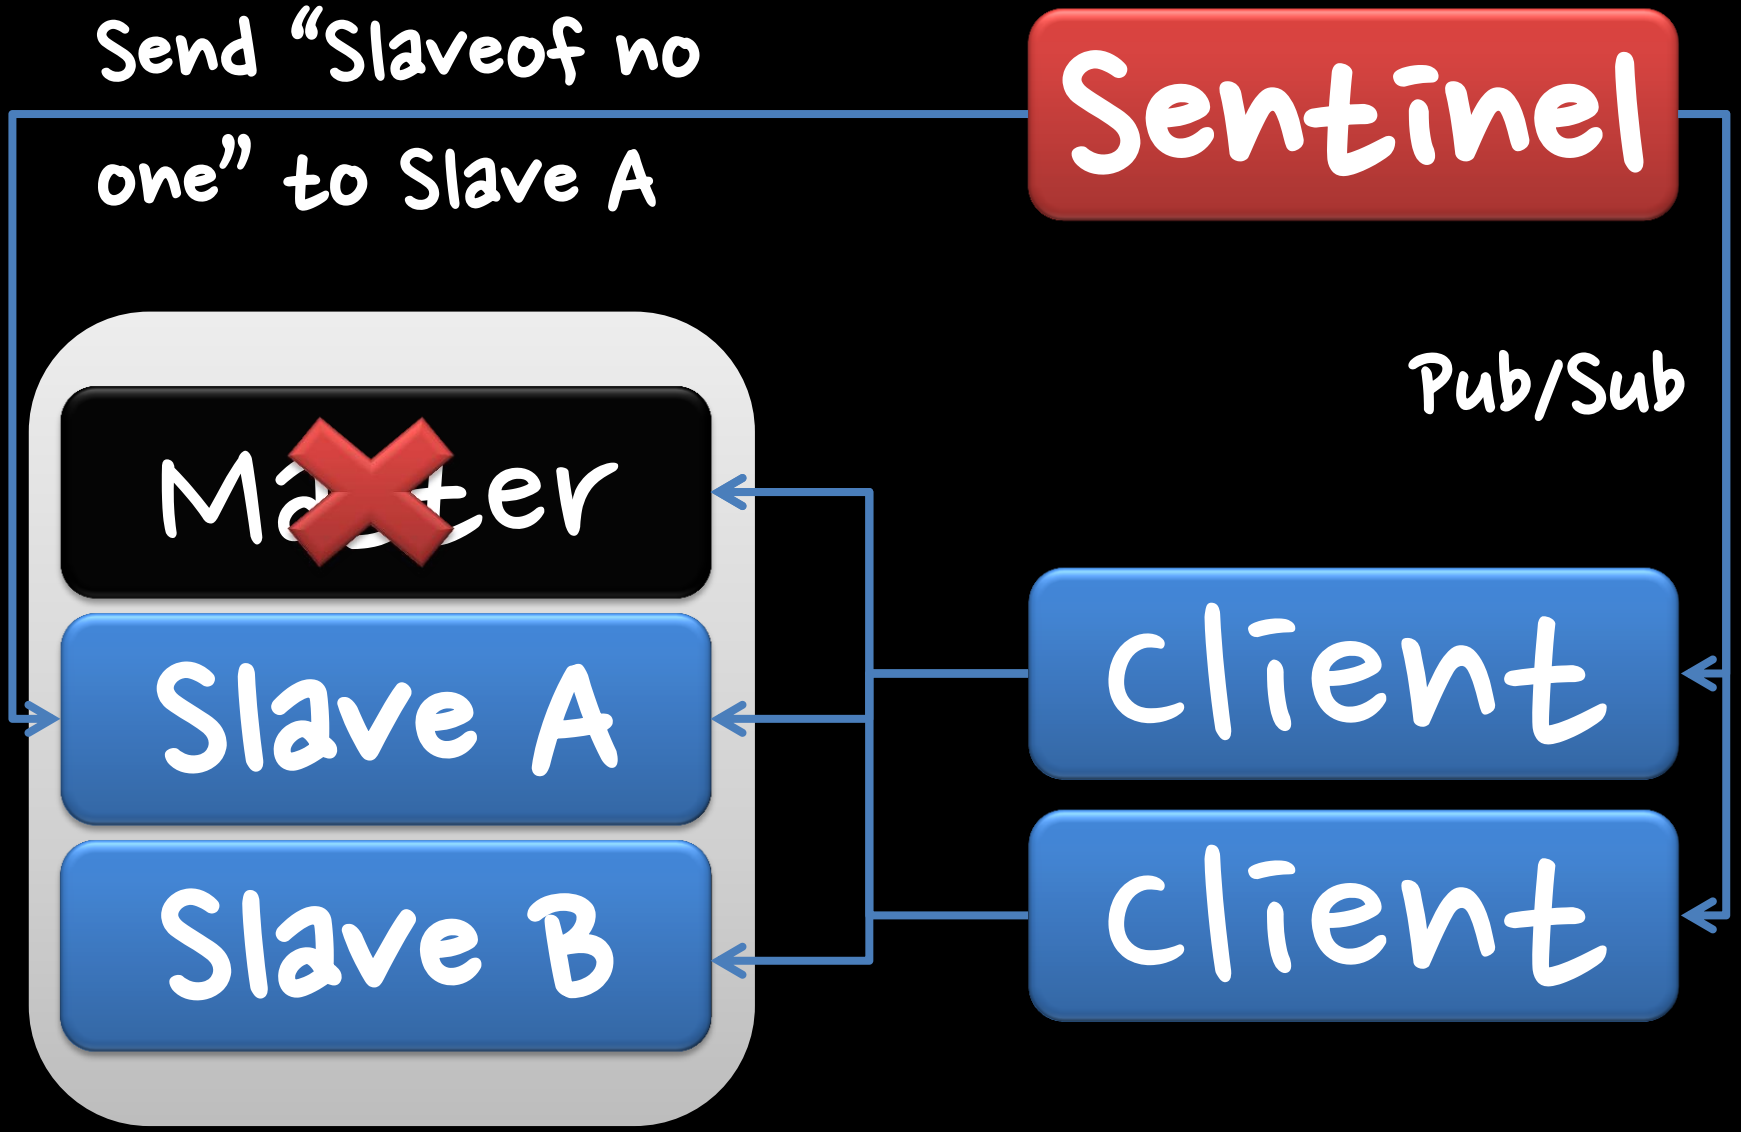
Send "Slaveof no one" to Slave A
Sentinel
Pub/Sub
Master
Slave A
Slave B
client
client

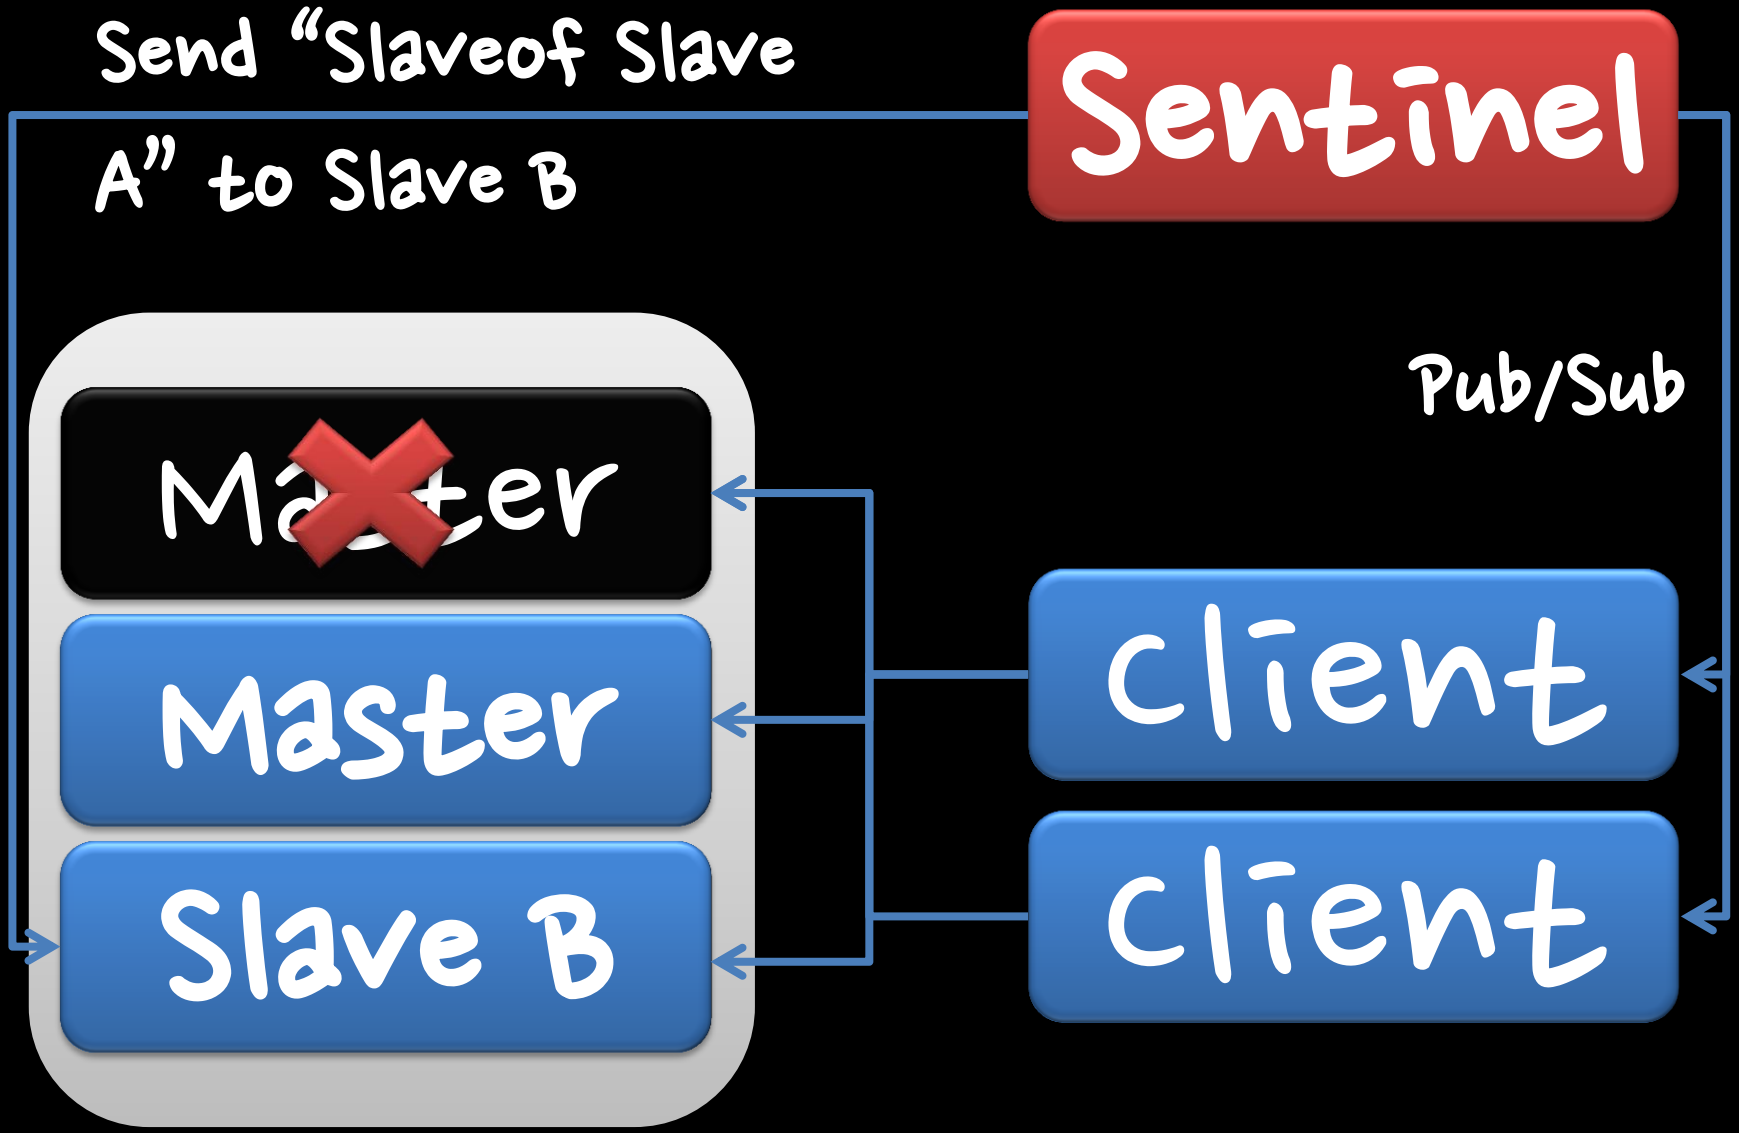
Send "Slaveof Slave A" to Slave B
Sentinel
Pub/Sub
Master
Master
Slave B
client
client

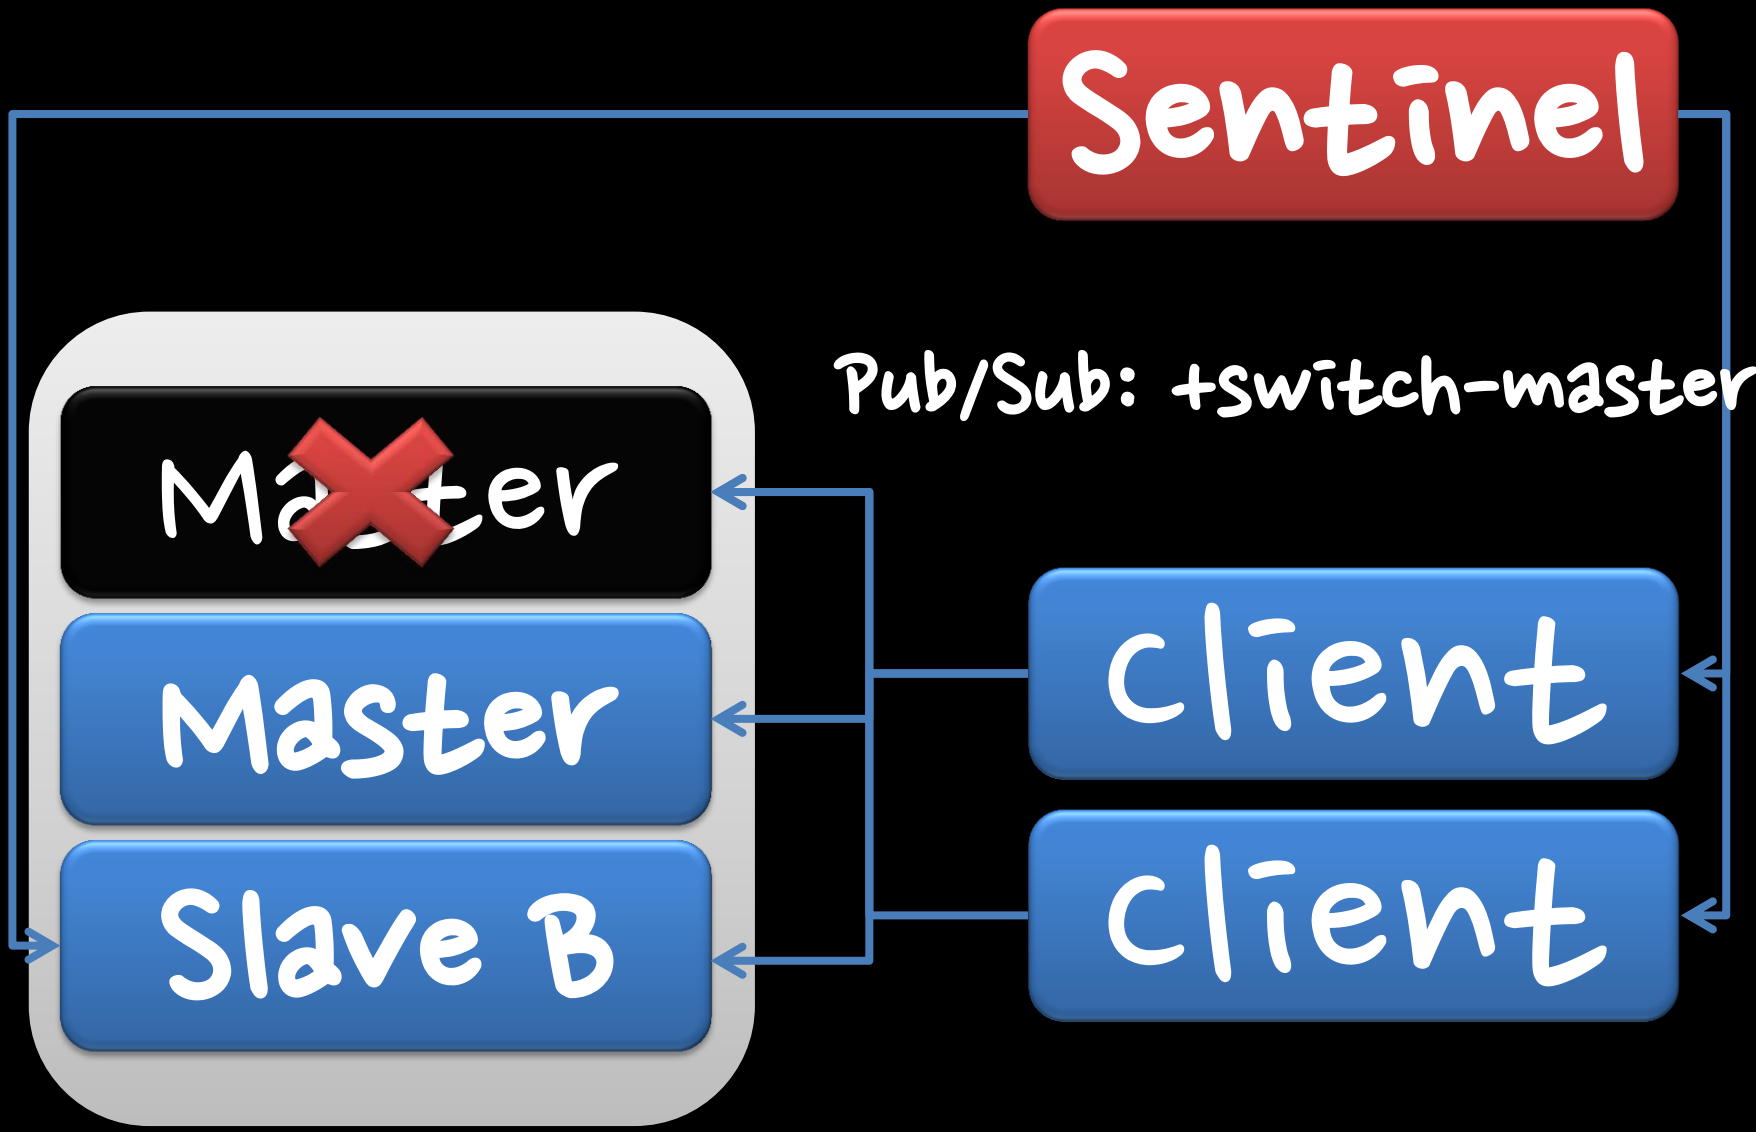
Sentinel
Pub/Sub: +switch-master
Master
Master
Slave B
client
client

Sentinel Internal

Sentinel Failover State Machine

State	내용	Next State
NONE		
WAIT_START		SELECT_SLAVE
SELECT_SLAVE		SEND_SLAVEOF
SEND_SLAVEOF		WAIT_PROMOTION
WAIT_PROMOTION		
RECONF_SLAVES		
DETECT_END		

Starting Point Of Sentinel

- sentinelTimer in sentinel.c
  - —called every 100ms
  - —sentinelHandleRedisInstance

SentinelRedisInstance

- master
  - If it is master, it is null.
- slaves
- sentinels

sentinelHandleDictOfRedisInstances

```
while((de = dictNext(di)) != NULL) {
        sentinelRedisInstance *ri = dictGetVal(de);

        sentinelHandleRedisInstance(ri);
        if (ri->flags & SRI_MASTER) {
            sentinelHandleDictOfRedisInstances(ri->slaves);
            sentinelHandleDictOfRedisInstances(ri->sentinels);
            ......
        }
}
```

sentinelHandleRedisInstance

- Reconnect to Instances
- Ping
- Asking Master State to other Sentinel
- check SDOWN
- check ODOWN

How to check Subjective Down

- When Sentinel get Ping Reply
  - Set last_avail_time as current
- Check
  - e = current -last_avail_time
  - e > last_avail_time
    - Subjective Down.

How to check Objective Down

- Ask to other sentinels is it down.
  - —sentinelAskMasterStateToOtherSentinels
- If other sentinels reply it is down.
  - —count them, if it is bigger than Quorum.
  - —It is objective Down.

How to find Redis

INFO Command

```
# Replication
role:master
connected_slaves:1
slave0:ip=127.0.0.1,port=6380,state=online,offset=1,lag=0
master_repl_offset:1
repl_backlog_active:1
repl_backlog_size:1048576
repl_backlog_first_byte_offset:2
repl_backlog_histlen:0
```

INFO Command

```
# Replication
role:slave
master_host:127.0.0.1
master_port:6379
```

+redirect-master

How to find Sentinel

Find Sentinel

Subscribe SENTINEL_HELLO_CHANNEL

In Master

Publish Sentinel Information

In sentinelPingInstance

Add New Sentinel

In sentinelReceiveHelloMessages

Using Pub/sub

```
*4
$8
pmessage
$1
*
$18
__sentinel__:hello
$58
127.0.0.1:26379:7d717f945afde99e6f82f825de052f17cab7e6f3:1
```

Is Sentinel Really Good?

Not Mature #1

Sentinel Conf

```
port 26379

sentinel monitor mymaster 127.0.0.1 6379 2
sentinel down-after-milliseconds mymaster 2000
sentinel can-failover mymaster yes
sentinel parallel-syncs mymaster 1
sentinel failover-timeout mymaster 900000
```

Quorum Count

```
sentinel monitor mymaster 127.0.0.1 6379 2
```

We should check this count.

Sentinel Conf

sentinel down-after-milliseconds mymaster 2000

If this value is Too small,
It can cause some trouble.
-sdown/+sdown loop

Compare with Zookeeper Conf

```
tickTime=2000
dataDir=/var/zookeeper
clientPort=2181
initLimit=5
syncLimit=2
server.1=zoo1:2888:3888
server.2=zoo2:2888:3888
server.3=zoo3:2888:3888
```

Not Mature #2

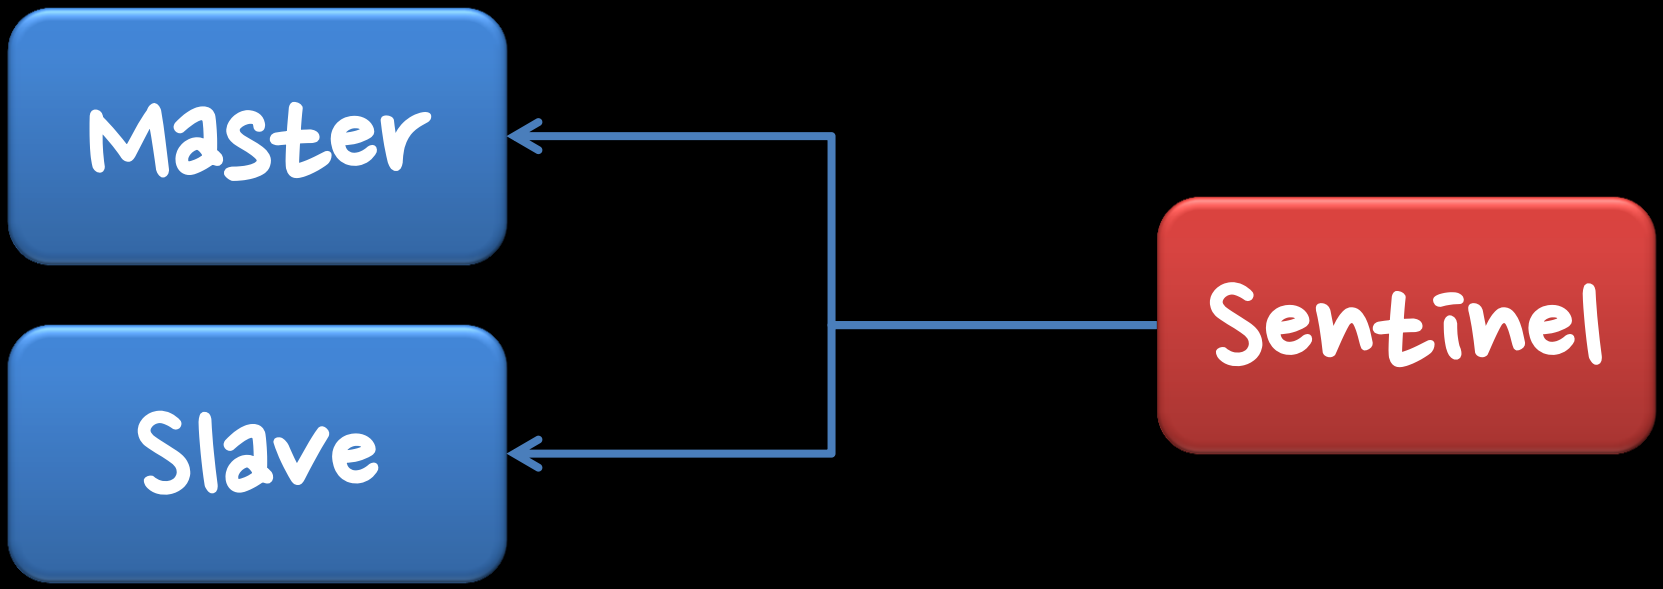
Master
Slave
Sentinel

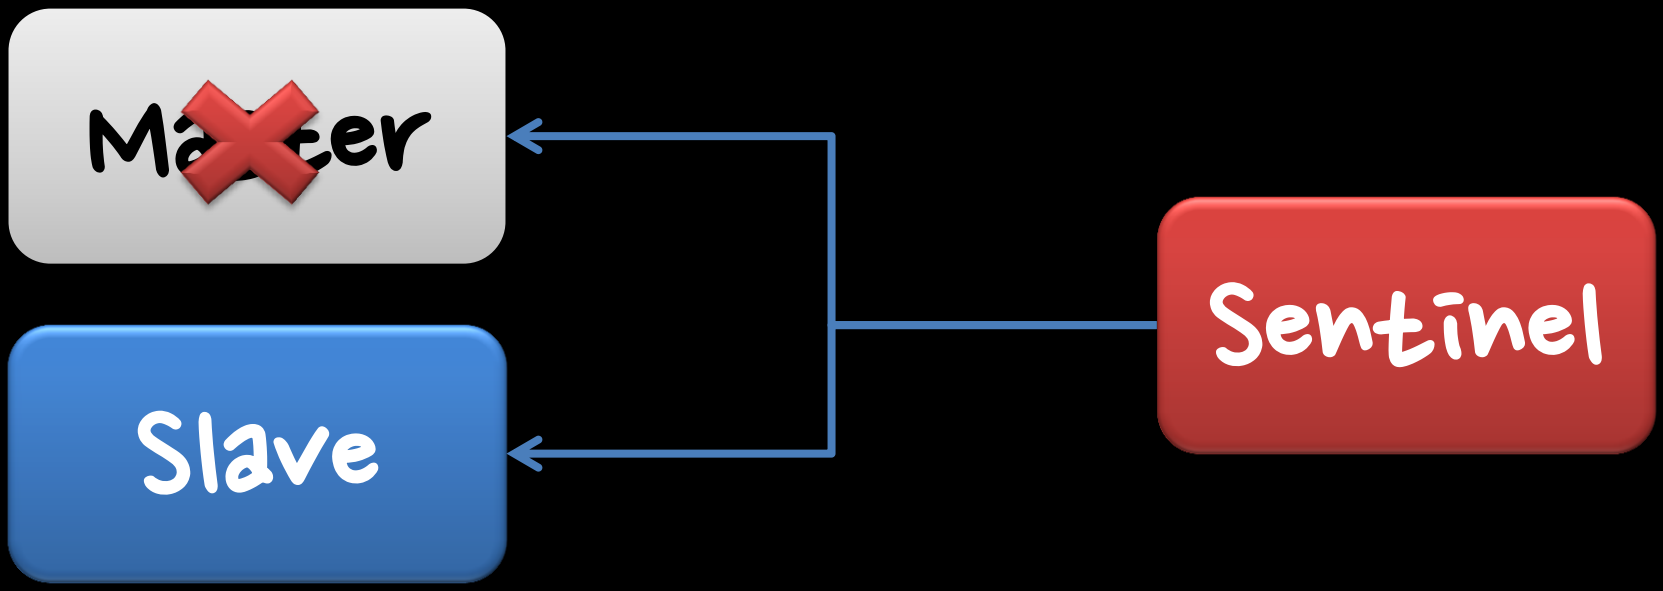
Master
Slave
Sentinel

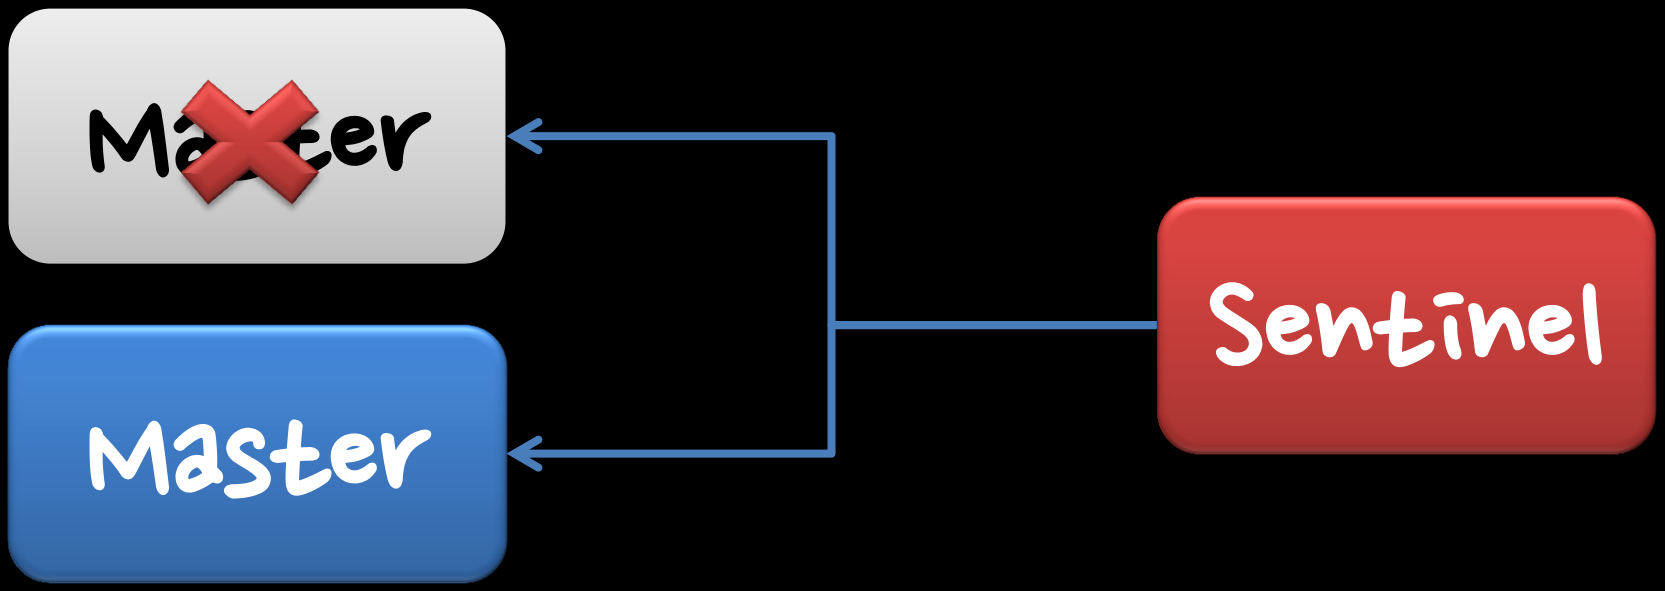
Master
Master
Sentinel

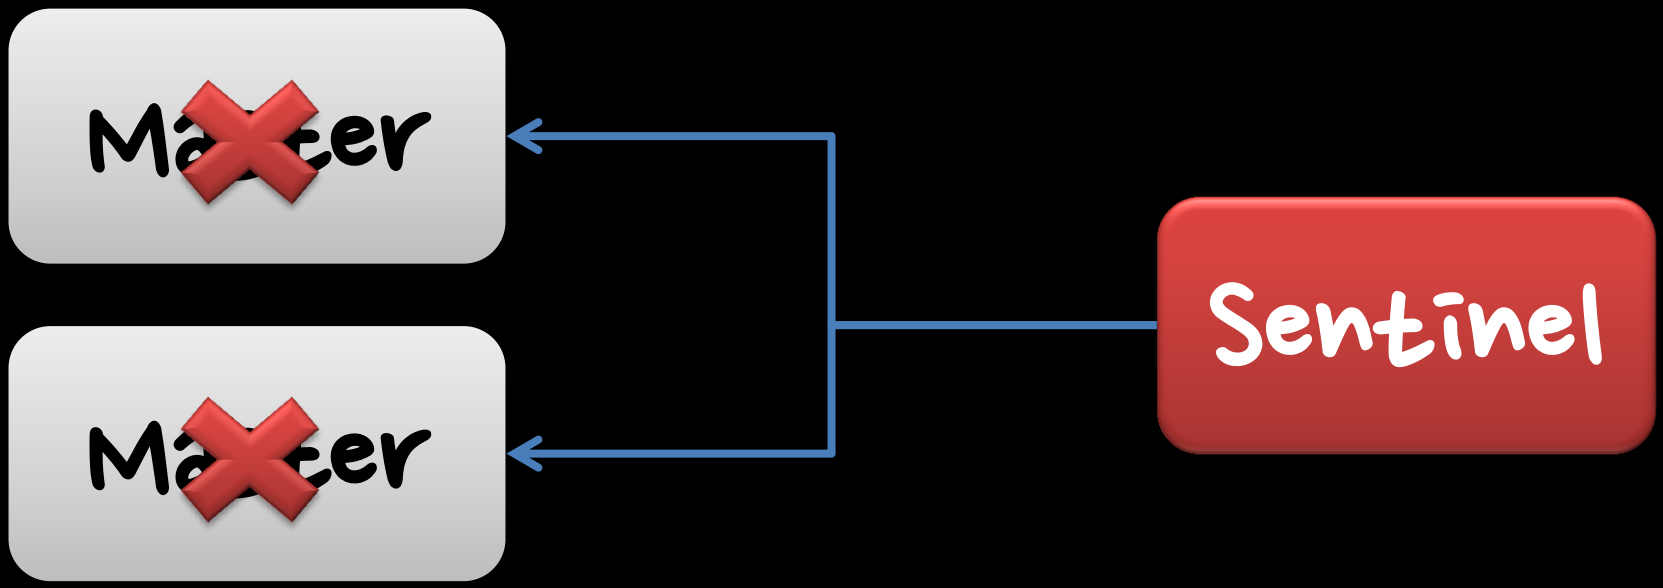
Master
Master
Sentinel

case #1

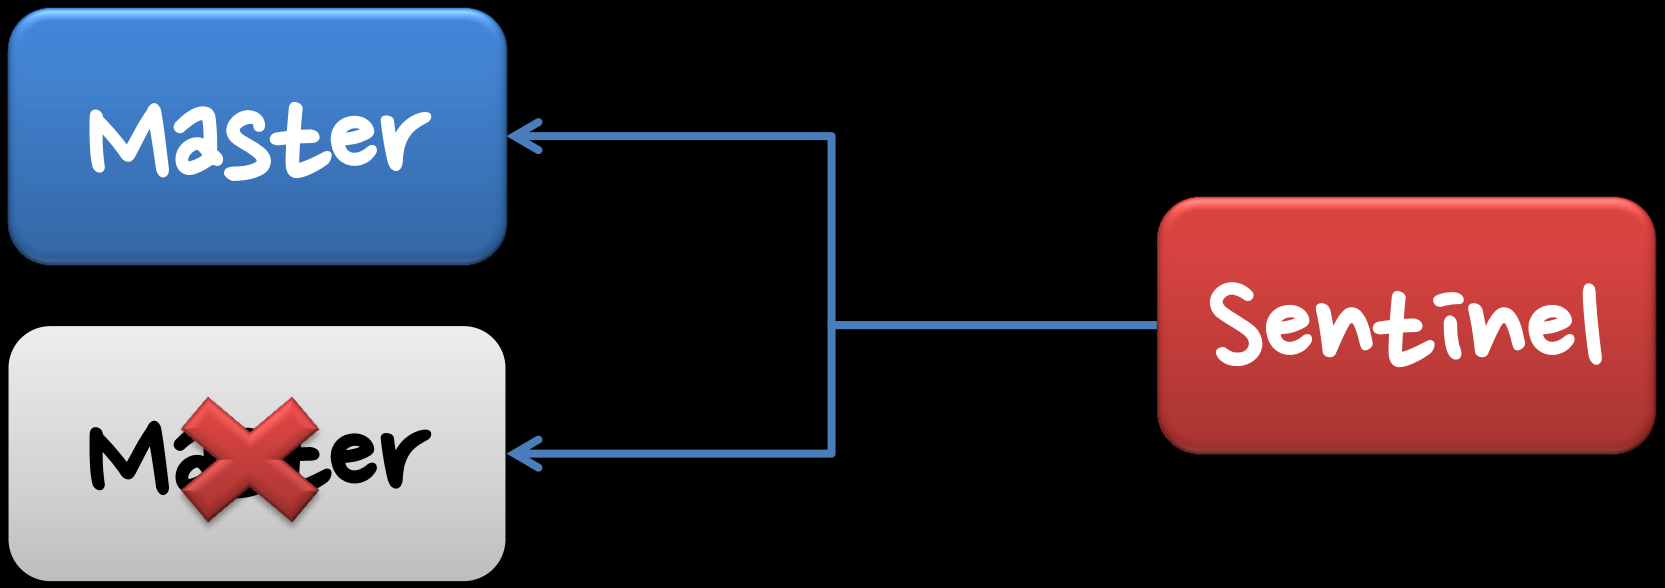

It's ok.

case #1

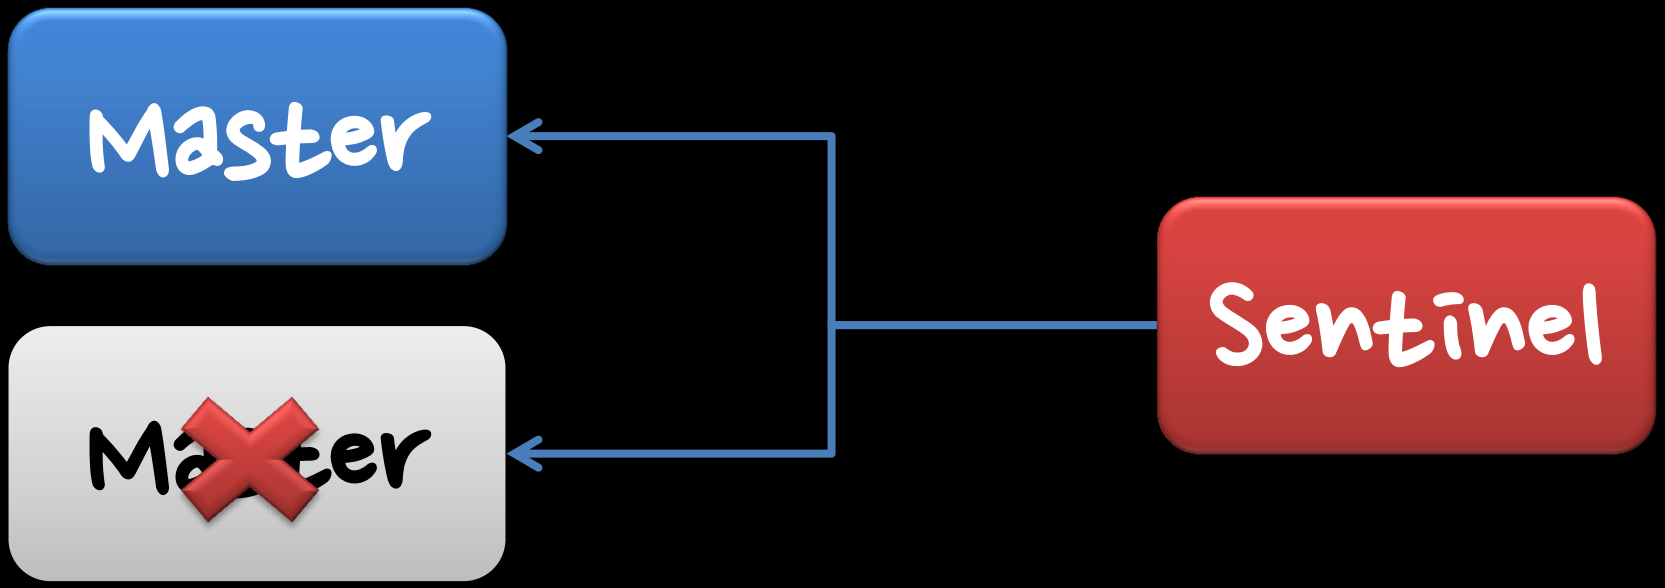

case #2

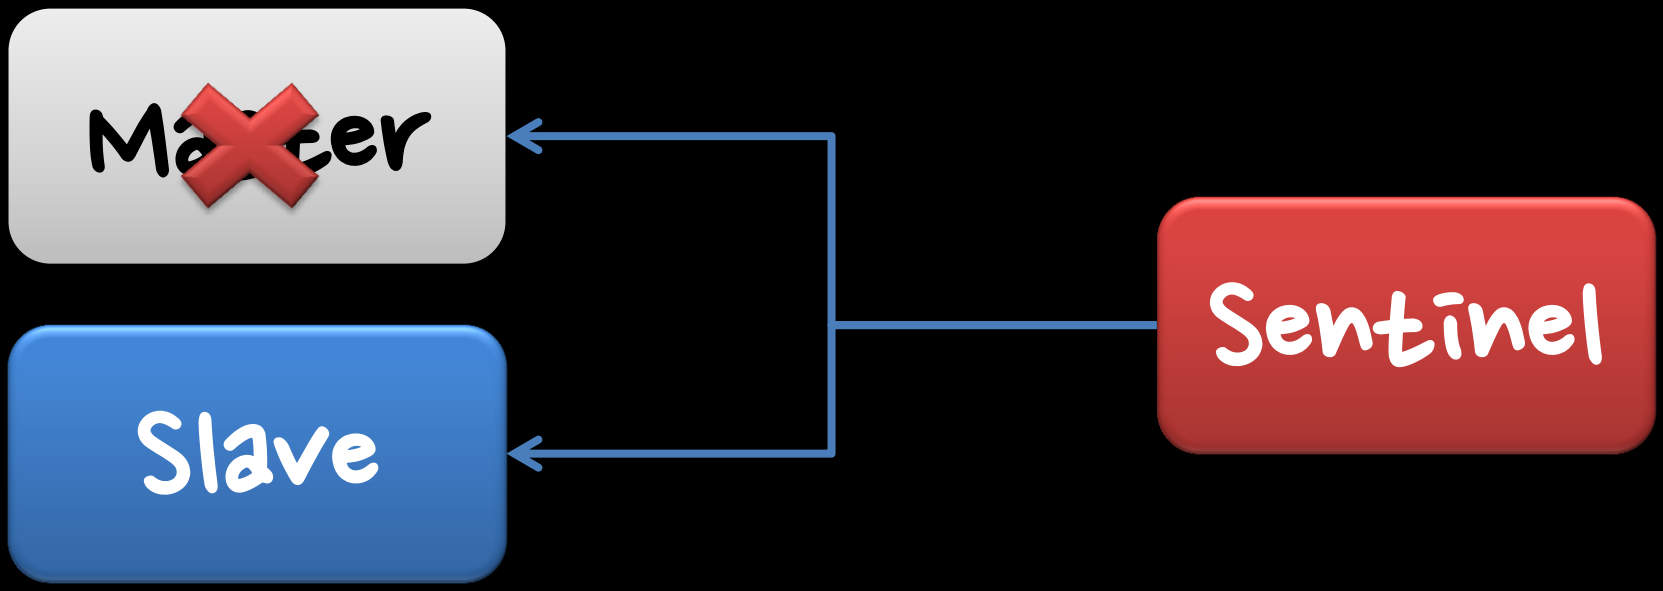

case #2

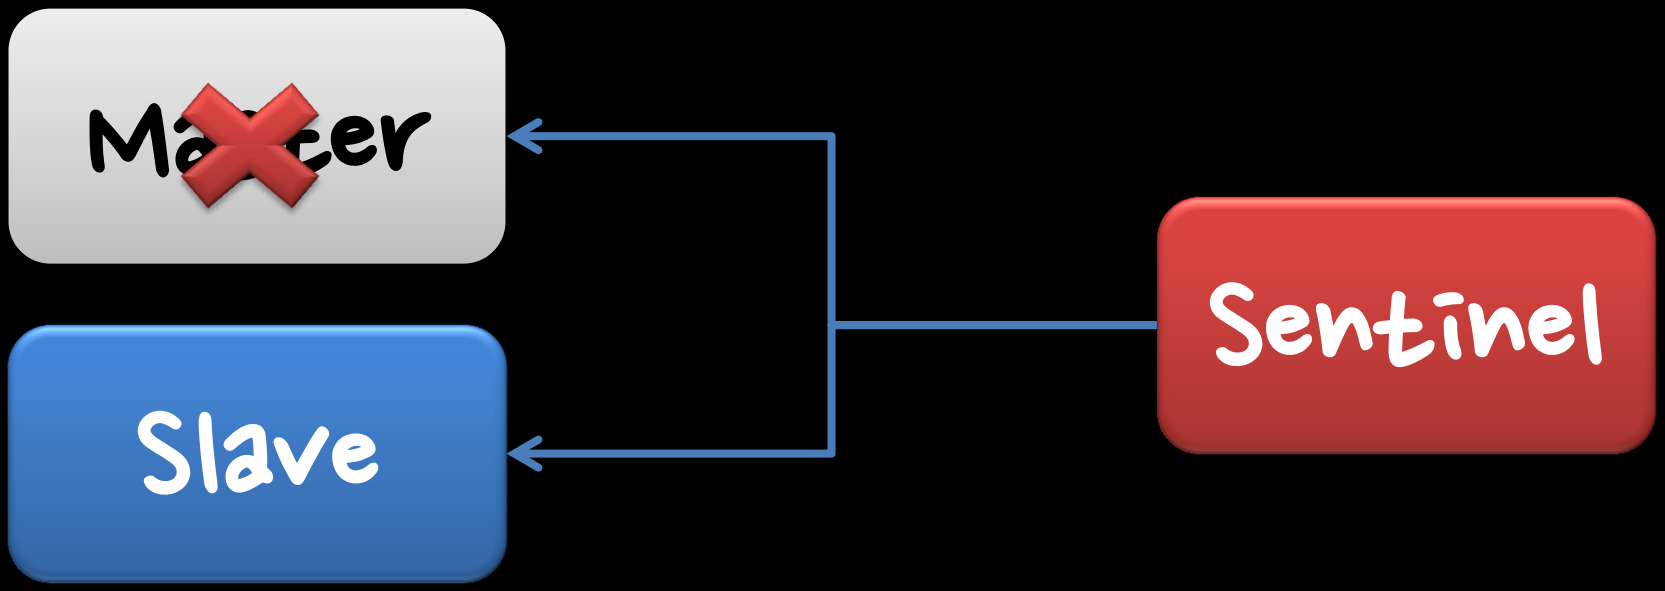

Sentinel can't promote

Slave to Master

why?

Sentinel Mechanism.

When master is changed to Slave,

Reconf with new Master.

Sentinel Mechanism.

When master is changed to Slave,
Reset information with new Master.

But, if new master failed, sentinel will not get any information from it.

Sentinel Mechanism.

```
if ((ri->flags & SRI_MASTER) && role == SRI_SLAVE && ri->slave_master_host)
    {
        sentinelEvent(REDIS_WARNING,"+redirect-to-master",ri,
            "%s %s %d %s %d",
            ri->name, ri->addr->ip, ri->addr->port,
            ri->slave_master_host, ri->slave_master_port);
        sentinelResetMasterAndchangeAddress(ri,ri->slave_master_host,
                                            ri->slave_master_port);
        return; /* Don't process anything after this event. */
    }
```

Not Mature #3

If redis is down, when redis is loading data(RDB, AOF), sentinel can't register redis.

conclusion

Even though, Sentinel is not mature.

It is useful.

Redis Sentinel Tips

1. Subscribe all Sentinels And ignore duplicated message

2. Sentinel has explicit failover command

sentinel mymaster failover

But can't select new master

Hacking is easy.

https://github.com/antirez/redis/pull/1126

3. Make Sentinel.conf same.

Sentinels can't share their status

Deploy conf to all sentinels with Script

And restart all sentinels.

4. If you want not to promote specific server to master

Set Slave-Priority = 0

Thank you!

Just one more thing

Twemprox

Redis

Redis-sentinel

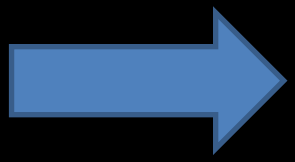

Sharded

Redis

Redis-Sentinel TwemProxy Agent

http://www.codeproject.com/Articles/656965/Redis-Sentinel-TwemProxy-Agent

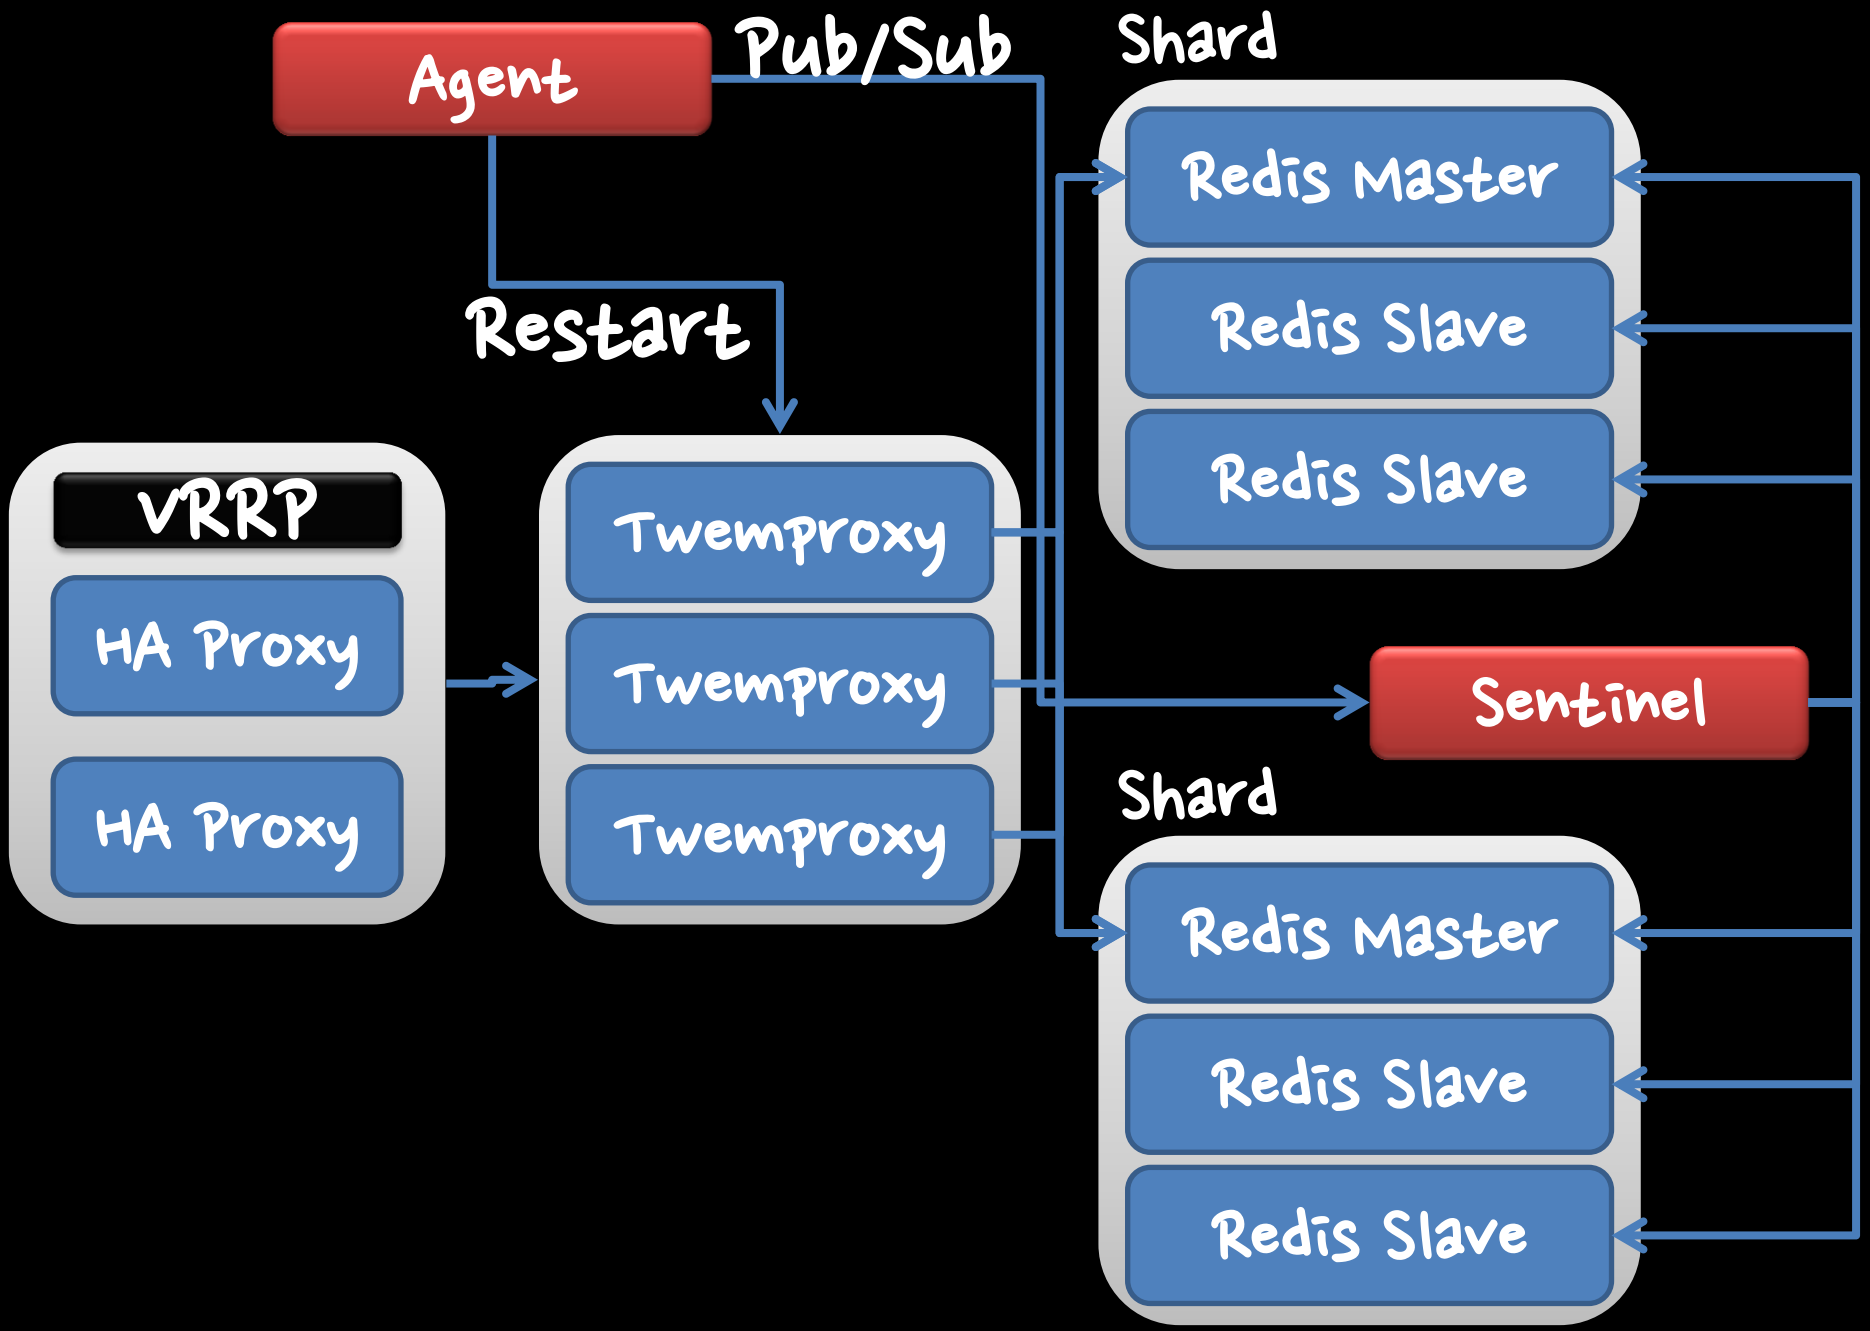
Agent
Pub/Sub
Shard
Redis Master
Redis Slave
Redis Slave
Restart
VRRP
HA Proxy
HA Proxy
Twemproxy
Twemproxy
Twemproxy
Sentinel
Shard
Redis Master
Redis Slave
Redis Slave