

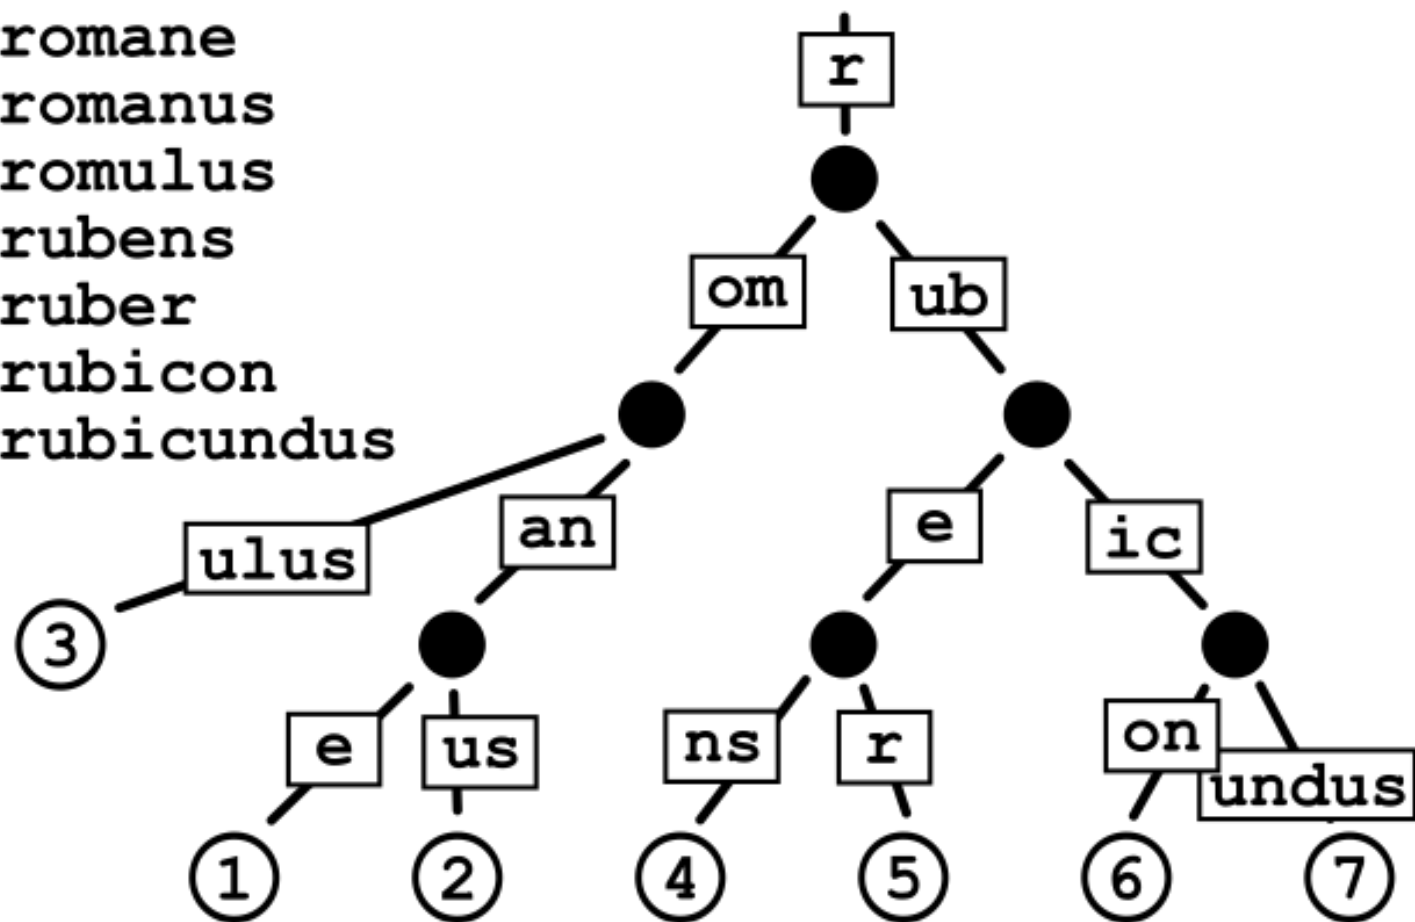
Rax, Listpack and safe contexts

@antirez

RAX

Redis radix tree implementation

- 1 romane
- 2 romanus
- 3 romulus
- 4 rubens
- 5 ruber
- 6 rubicon
- 7 rubicundus



HDR

a b c

a-ptr b-ptr c-ptr

val-ptr

Rax facts

- 2600 lines of code
- Less memory used
- Faster than dict.c for most use cases
- Ordered! Powerful iterators
- 100% fuzz tested

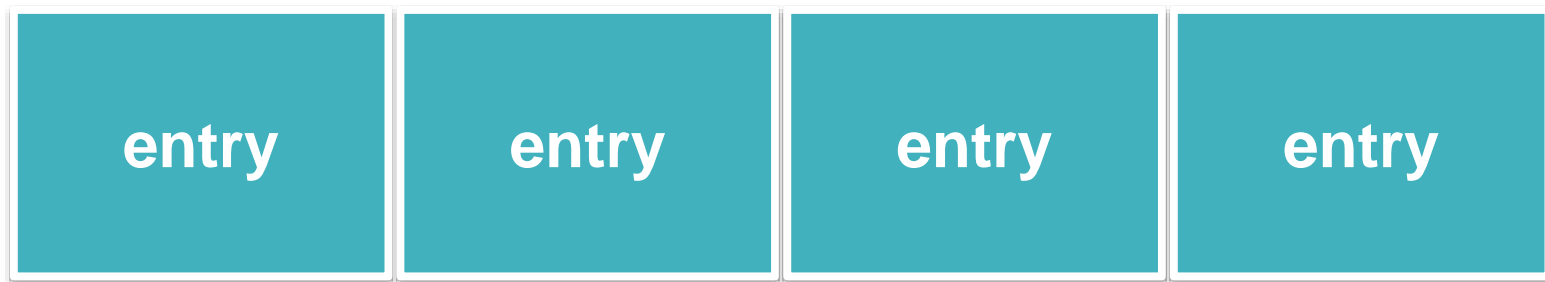
Rax usages

- Cluster key <-> slots tracking (done in 4.0)
- Streams data structure (4.2)
- Memory efficient Hashes and Sets (4.2)
- No longer double encoded! \o/
- Main dictionary? Expires?

Listpack



Main problem...



Ziplist encoding



Listpack encoding



Reverse length



Listpack facts

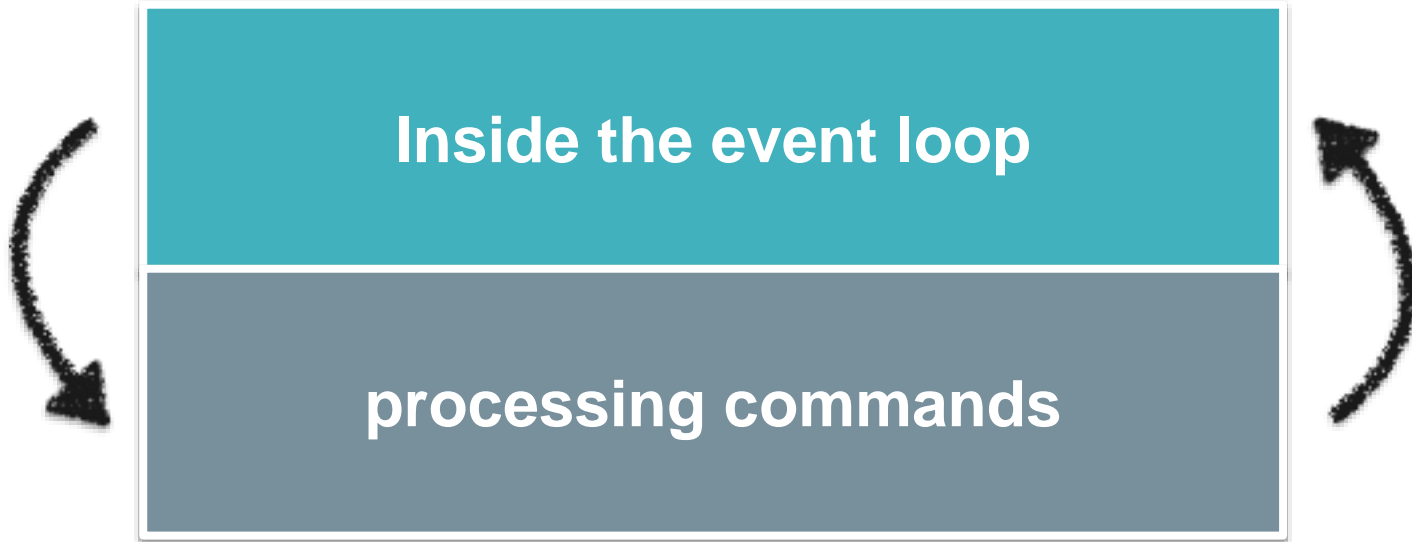
- We have a specification! By Yuval, Oran and me
- 1/3 the source code size of ziplist :-)
- More memory efficient for Redis data model
- work in progress (alpha published)
- TODO: 100% fuzz testing

Thread safe contexts

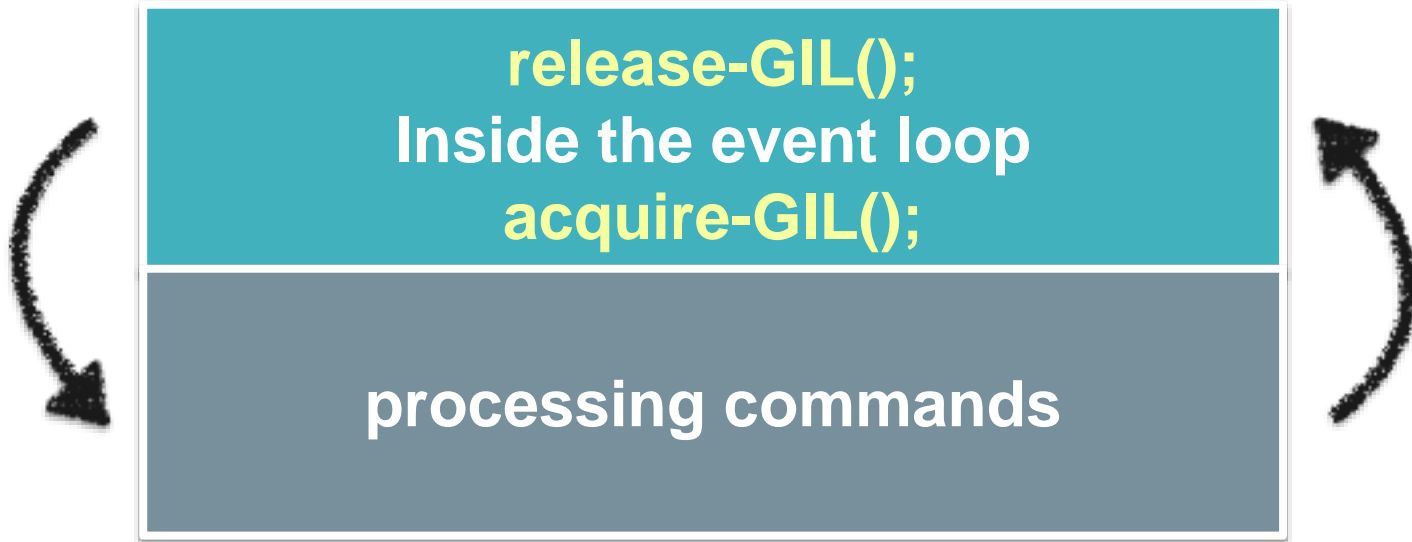
TSC: what we had

- blocking commands in modules
- RedisModule_UnblockClient() from thread
- No Redis Modules API from threads :-(

Redis event loop



Redis event loop



Non blocking KEYS *

```
do {  
    RedisModule_ThreadSafeContextLock(ctx);  
    RedisModuleCallReply *reply =  
        RedisModule_Call(ctx,  
            "SCAN", "l", (long long) cursor);  
    RedisModule_ThreadSafeContextUnlock(ctx);  
  
    ...  
}
```

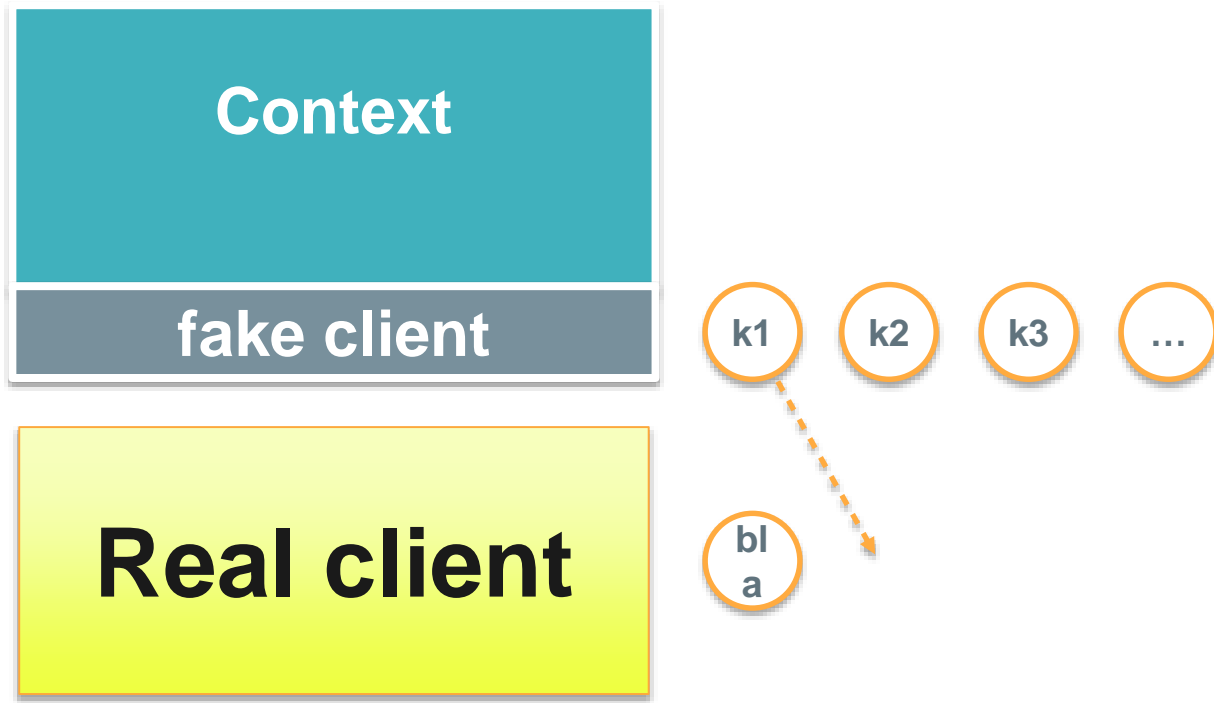
```
/* Note: no context lock to reply. */  
  
do {  
    ...  
    for (size_t j = 0; j < items; j++) {  
        RedisModuleCallReply *ele = ...;  
        RedisModule_ReplyWithCallReply(ctx,ele) ;  
    }  
}
```

Replies #1

Context

fake client

Replies #2



Replies #3



$O(1)$ operation



Future

Modules threading future

- Key-locking
- read-write locks
- Automatic blocking of clients accessing busy keys
- No need for data structures-level locks

Thank You

Salvatore Sanfilippo
Redis Labs

Tweet me at @antirez